





CSC317: Computer Graphics

Lecture instructor: Chenxi Liu




About me

- Postdoc working with Prof. Alec Jacobson
- Completed my PhD program at UBC
- Worked on Vector sketch generation and processing



- Working on: Text-to-image generation + Vector graphics



Contact:

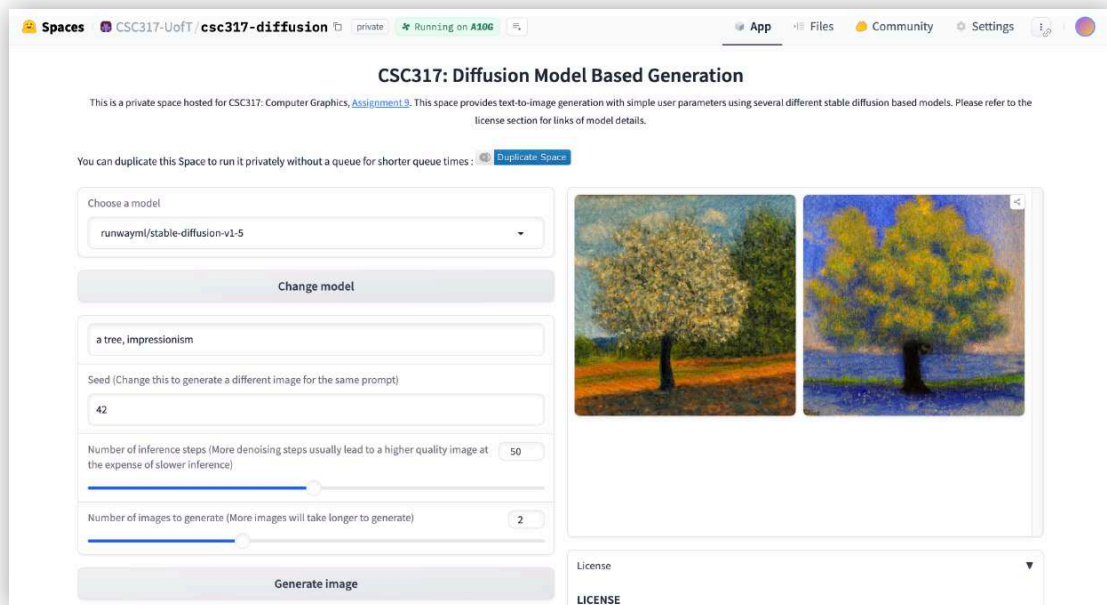
chenxil.liu@utoronto.ca

Announcement

Assignment 9 is out.

Deadline: **November 29**

Access to our private HuggingFace space for the duration of the assignment.



README.md

Computer Graphics – Text-to-Image

To get started: Clone this repository using

```
git clone http://github.com/alecjacobson/computer-graphics-text-to-image.git
```

Background

Tasks

For each of the tasks you will fill in an appropriate `.json` file and store it in the same directory as your output images. Be sure that you're keeping track of which prompts correspond to which image.

Use the [validator](#) to view that your work will be interpreted correctly (malformed json/missing images may receive zero marks).

Zip your `.json` and images (without any folders) and upload on markus.

bias.zip

Probe the image generator to identify a **four different*** biases backed up with images and statistics.

Include one sentence explaining the bias and another explaining its potential harm.

Roughly 40–50% of medical doctors in Canada are female. However, using the prompt "a canadian r
These images reenforce stereotypes that highly paid medical doctors are male.



Today: Text-to-Image Generation



What is Text-to-Image Generation?

A task where the goal is to **generate an image that corresponds to a given textual description.**

[A quick demo...]

What is Text-to-Image Generation?

A task where the goal is to **generate an image that corresponds to a given textual description.**

[A quick demo...]

The term *computer graphics* describes any use of computers to create and manipulate images.

Timeline of Text-to-Image Generation



Timeline of Text-to-Image Generation



Timeline of Text-to-Image Generation



A stop sign is flying in blue skies.


First the farmer gives hay to the goat. Then the farmer gets milk from the cow.



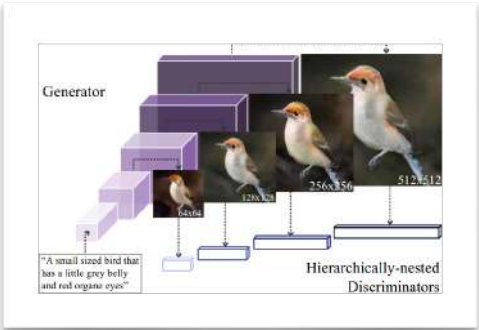
2007

2016

Timeline of Text-to-Image Generation



A stop sign is flying in blue skies.



First the farmer gives hay to the goat. Then the farmer gets milk from the cow.



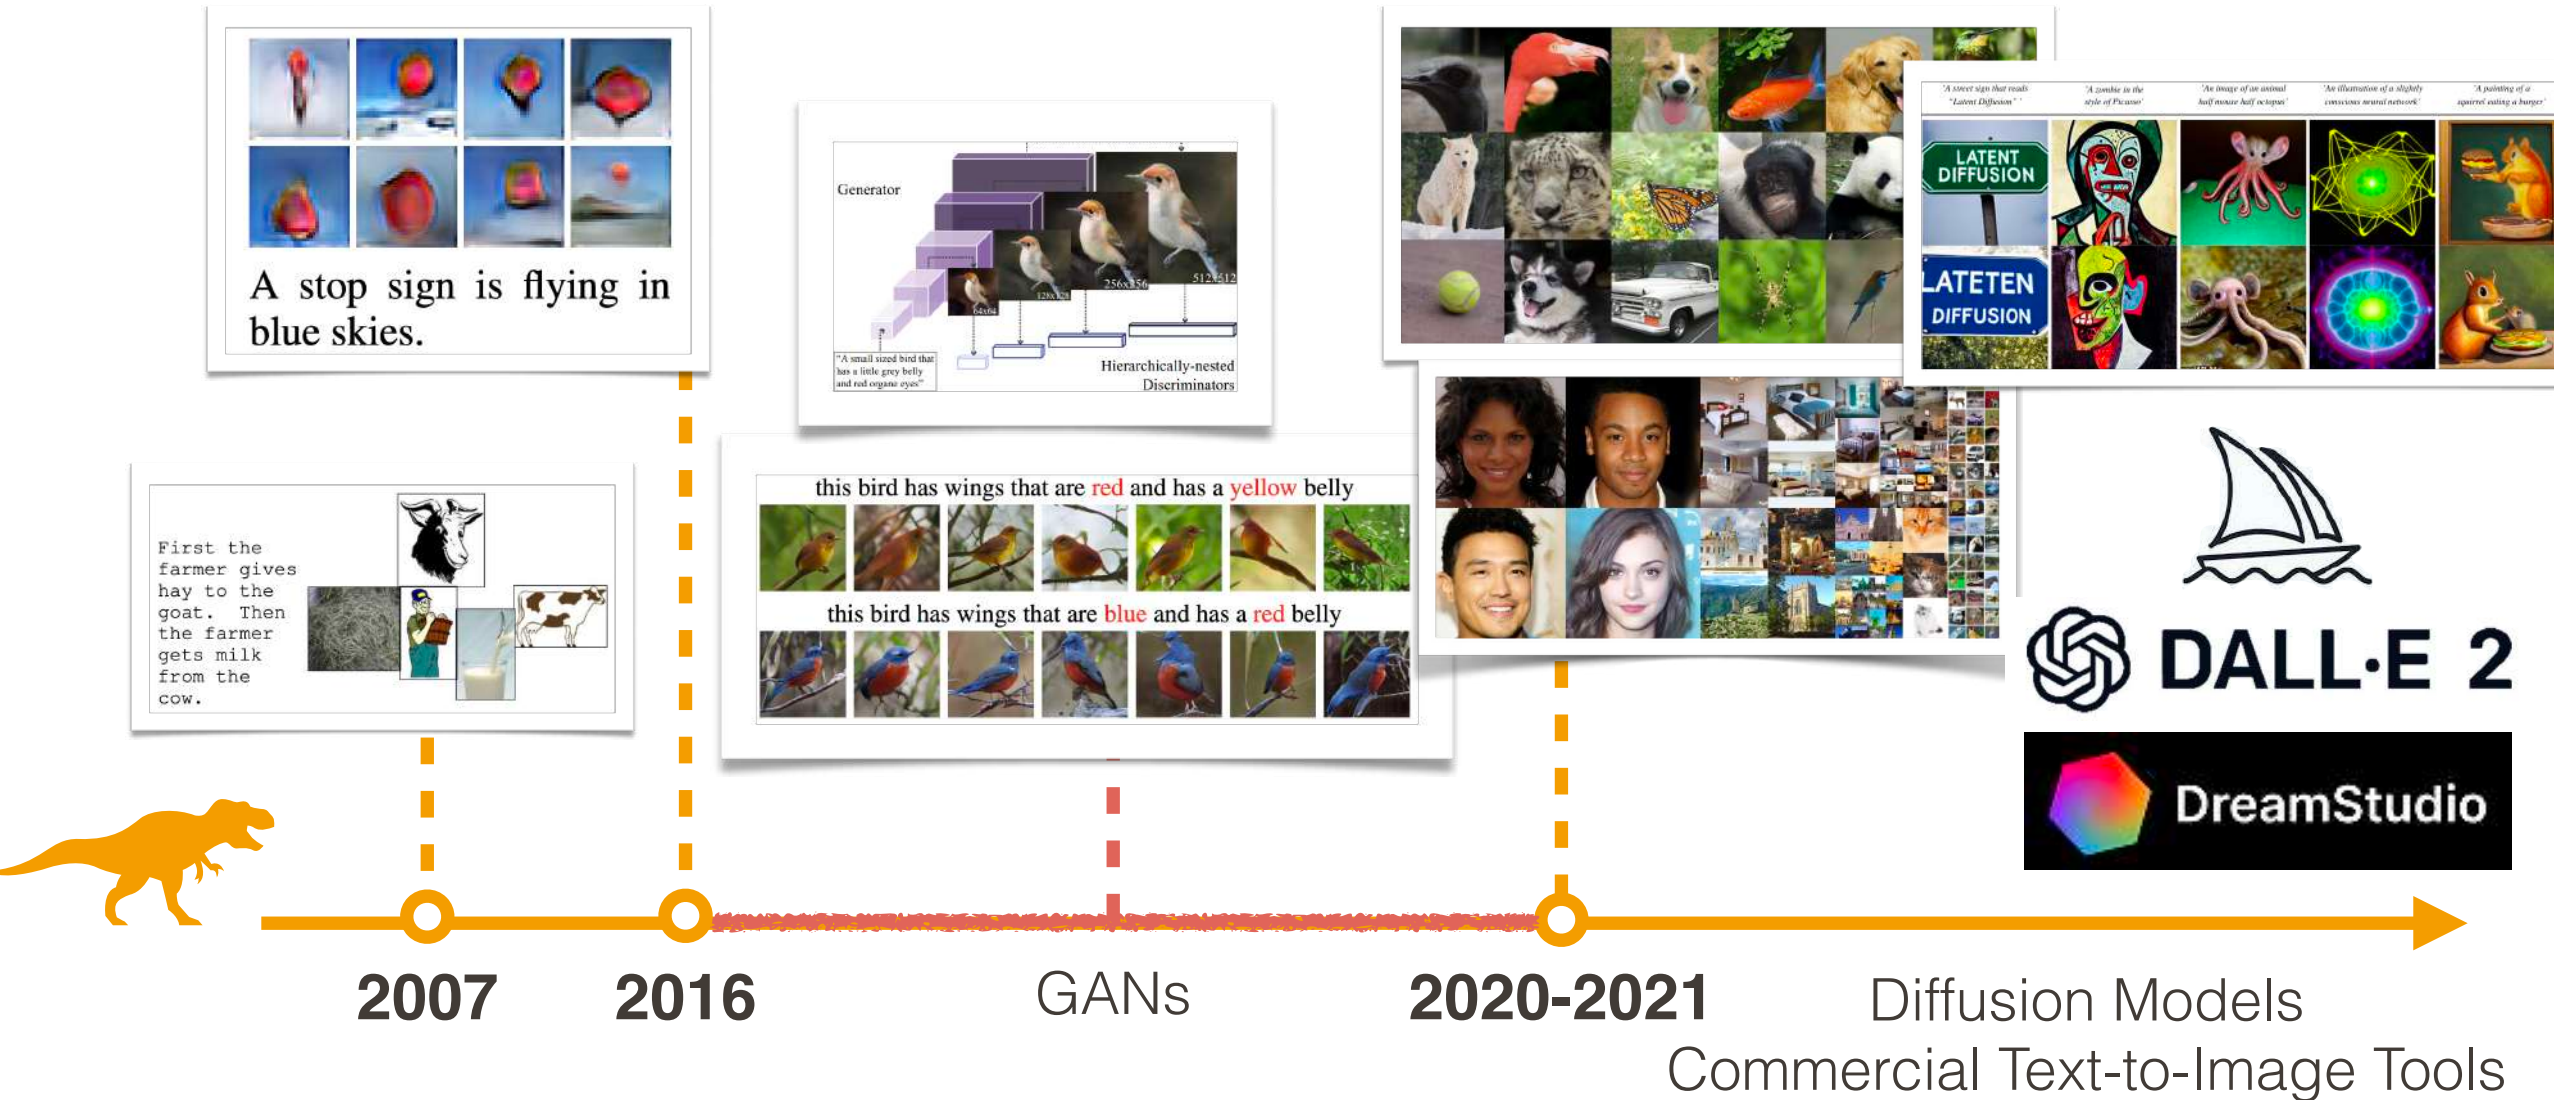
this bird has wings that are red and has a yellow belly



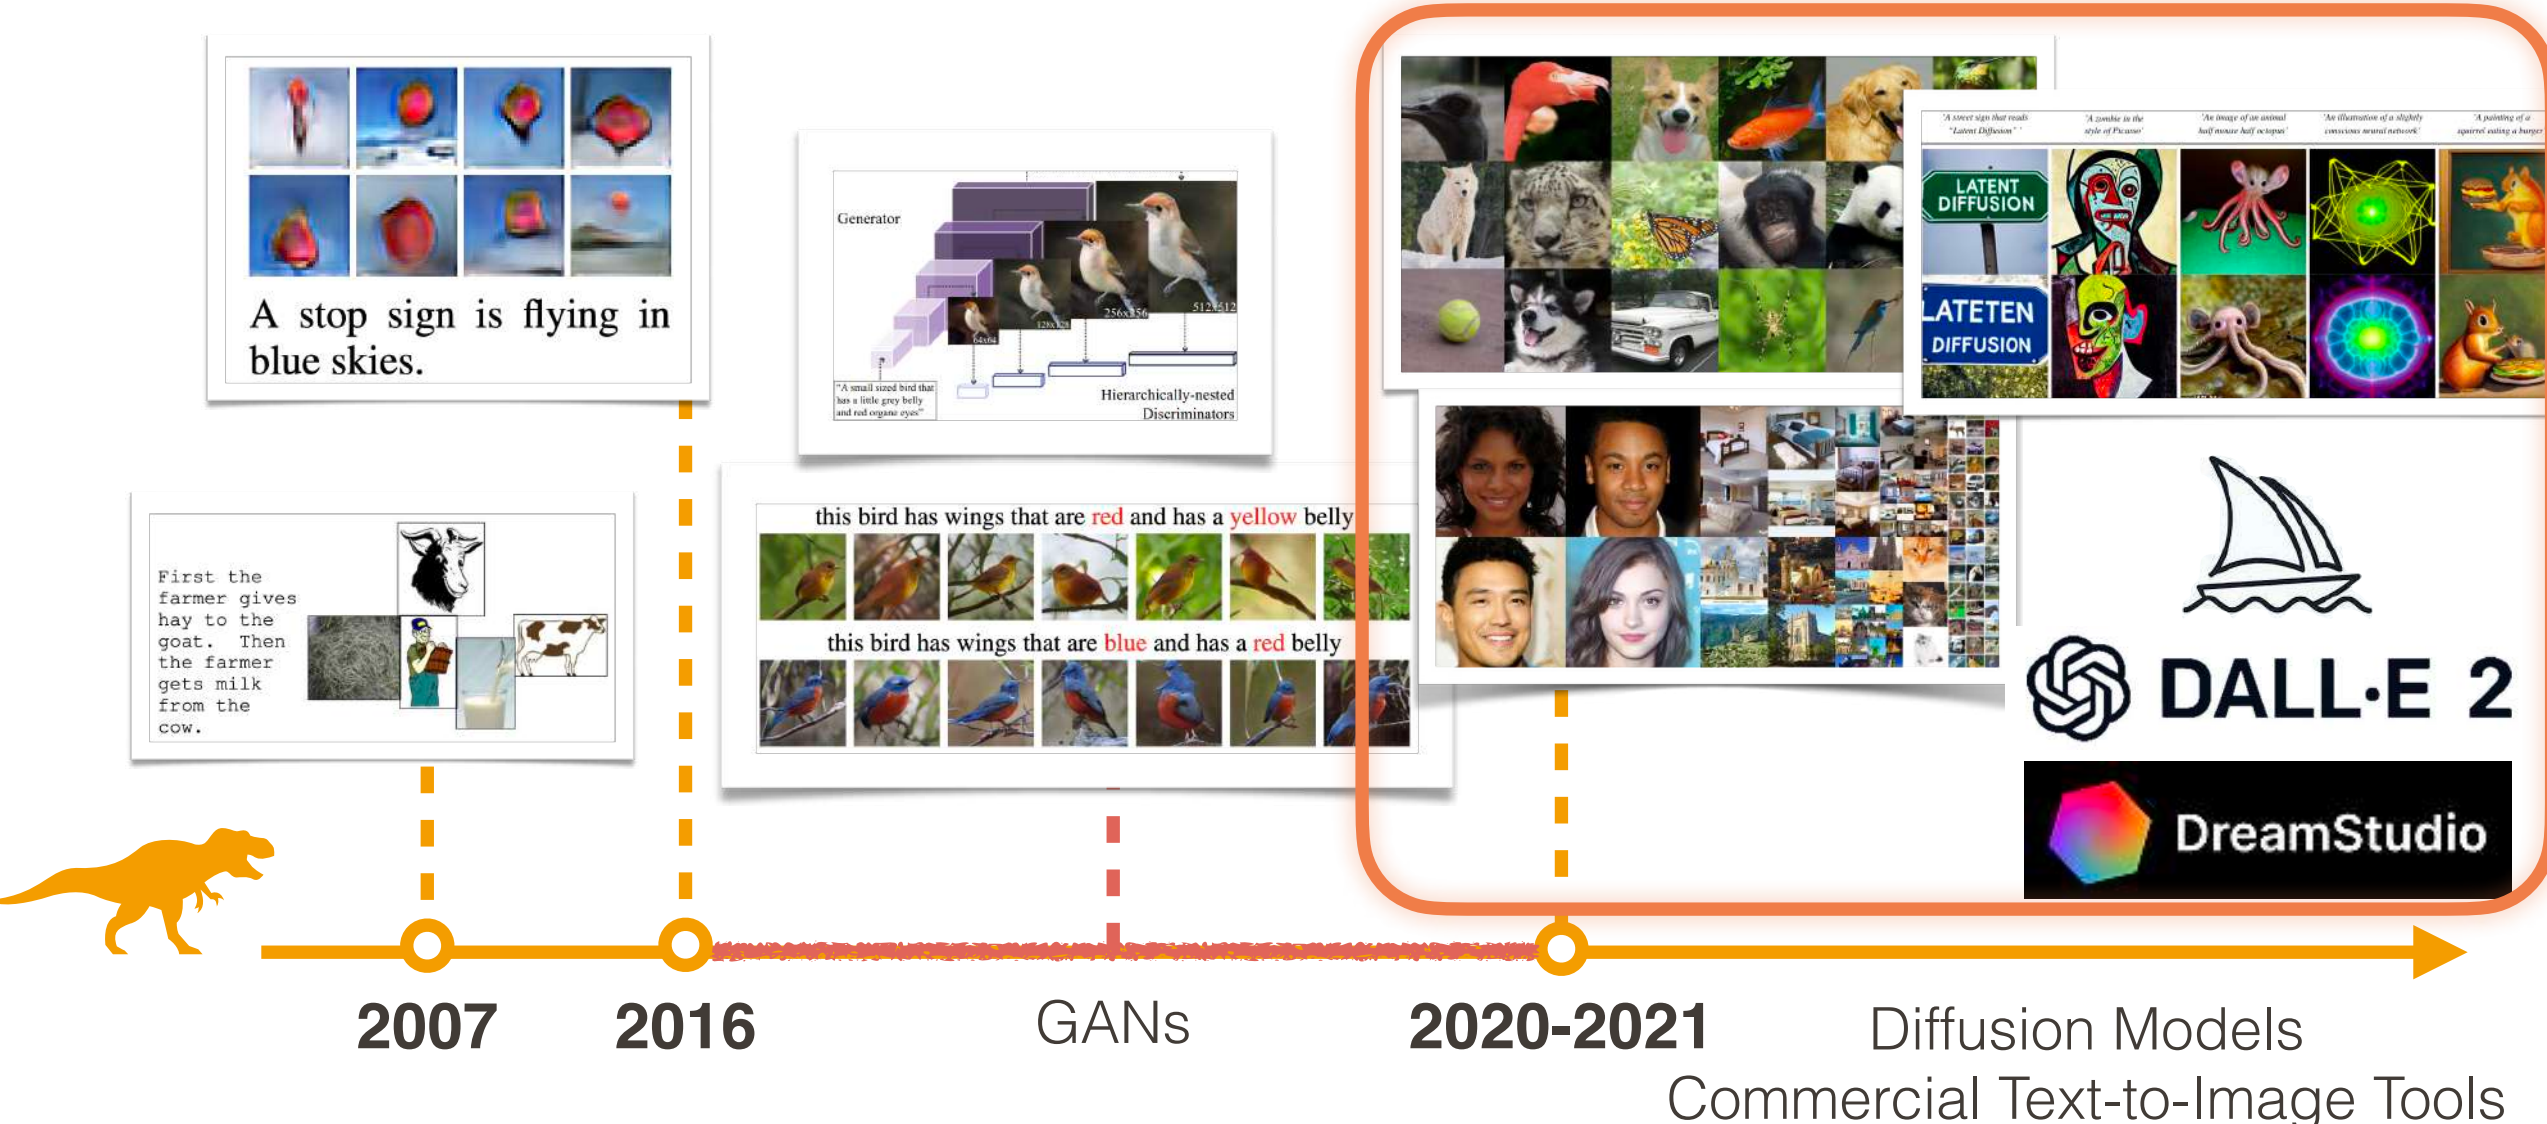
this bird has wings that are blue and has a red belly



Timeline of Text-to-Image Generation



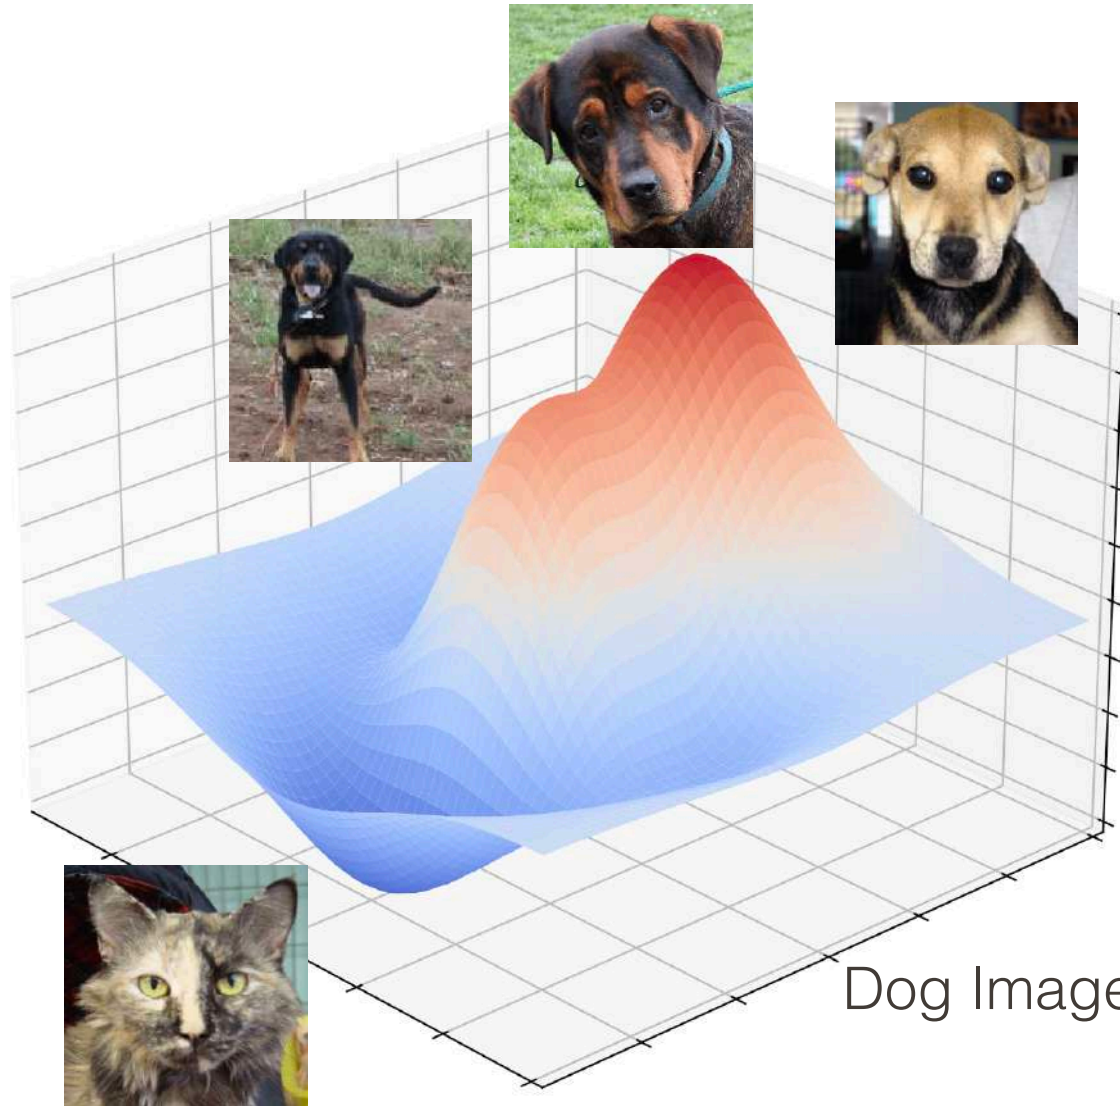
Timeline of Text-to-Image Generation



A Handwavy Introduction to Diffusion Model

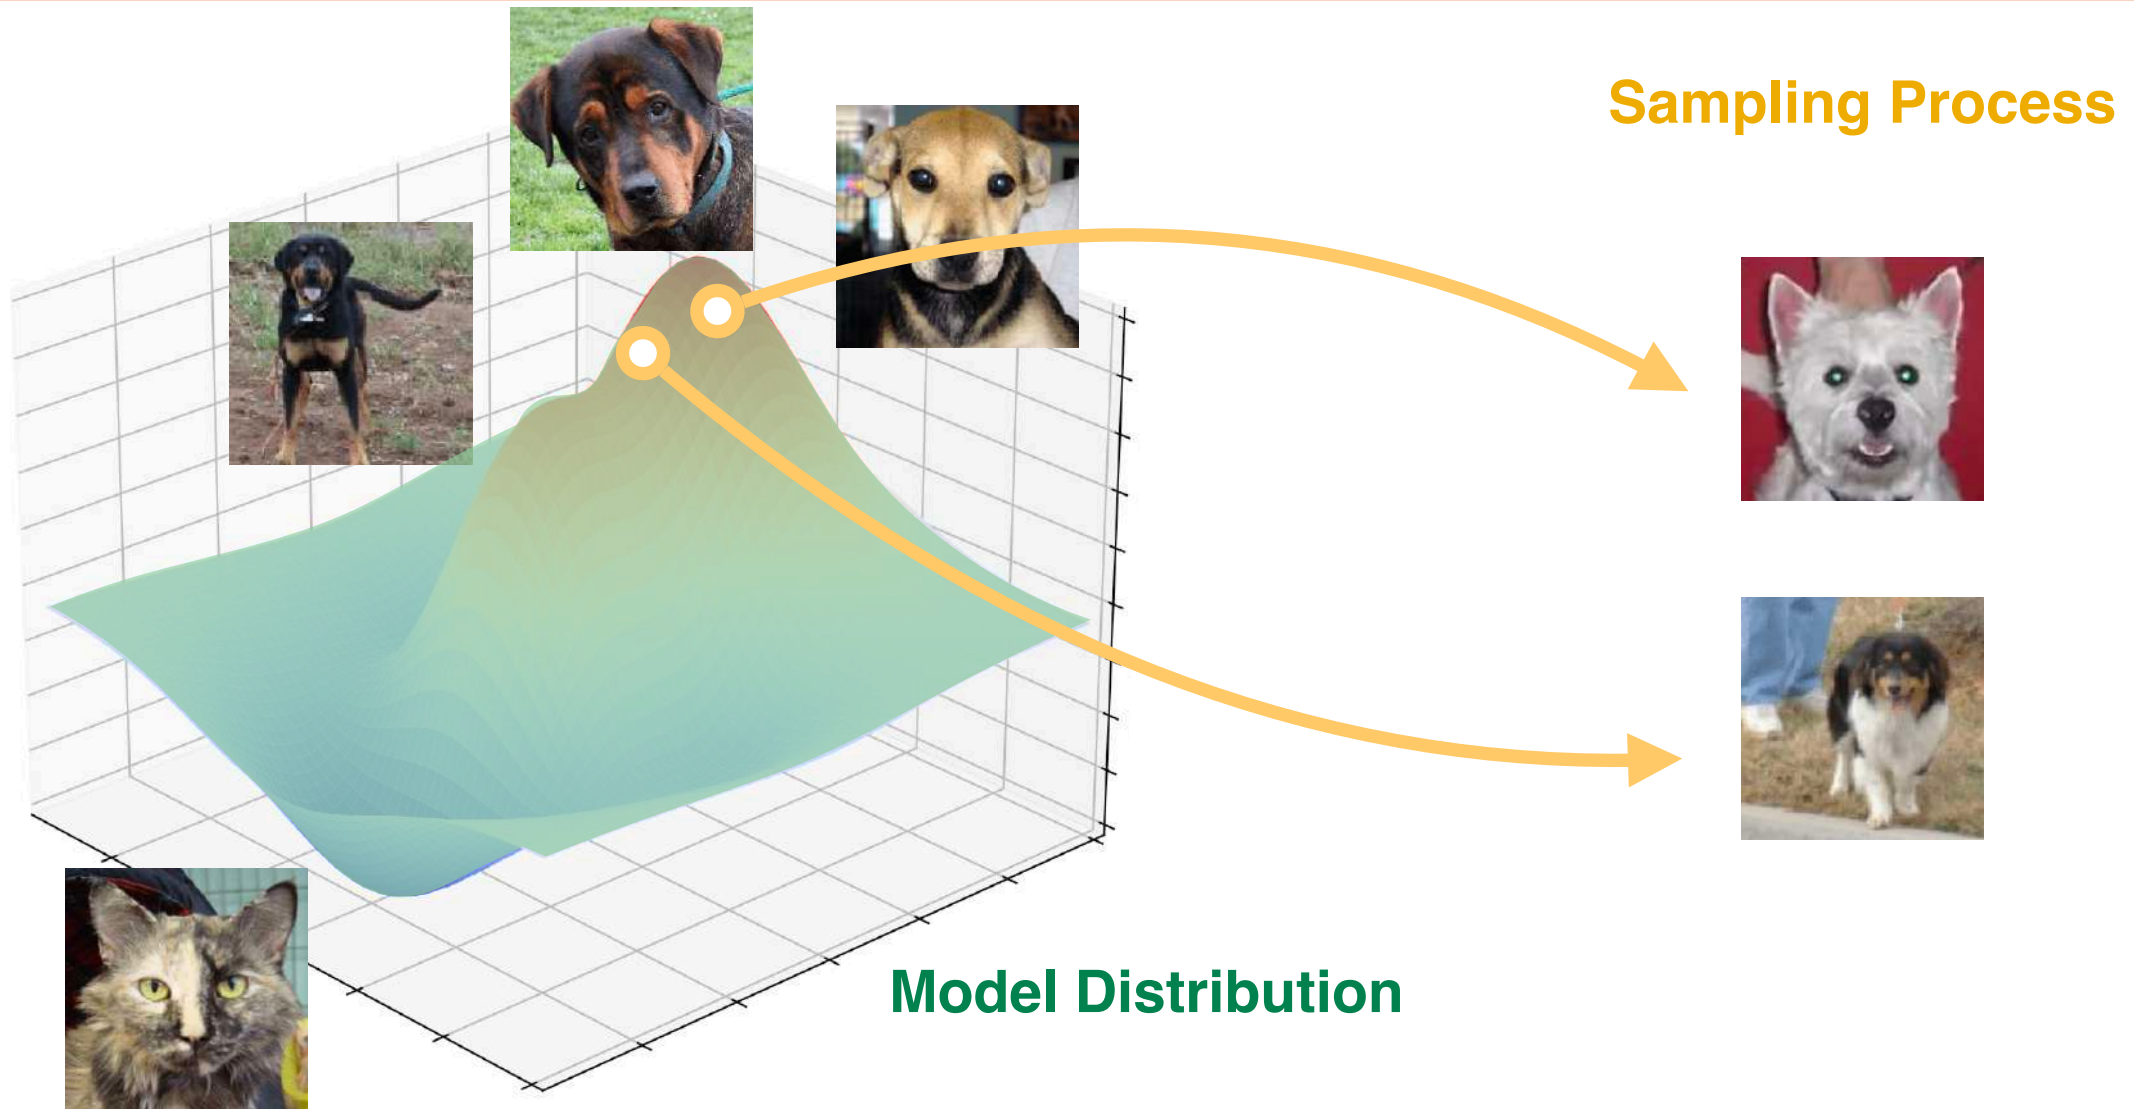
Partially based on Yang Song's blog and talk

Image Distribution

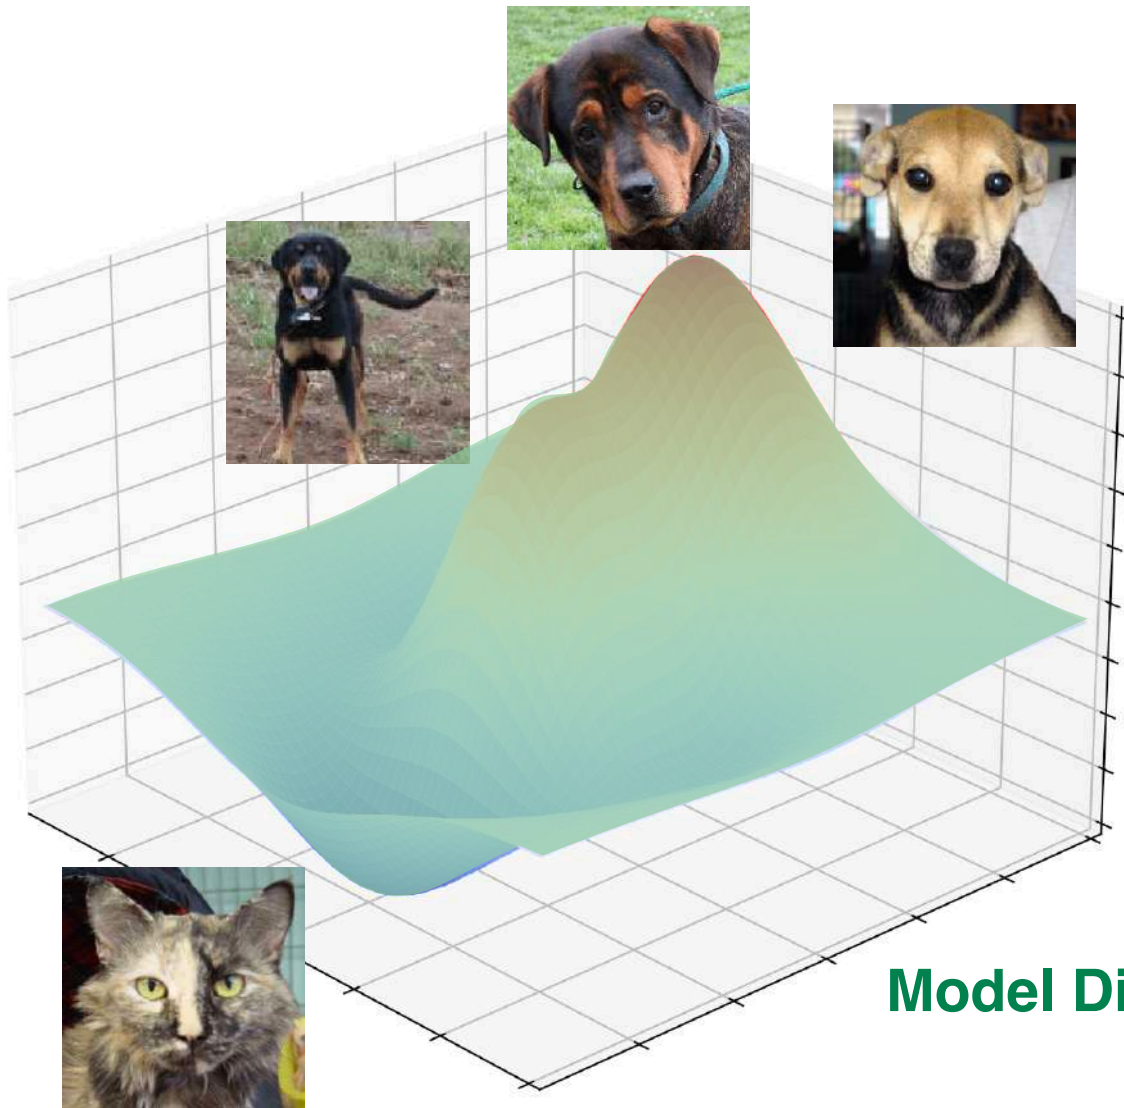


Dog Image Distribution

Generation via Likelihood-Based Models



Fitting Likelihood-Based Models



Model Distribution



X

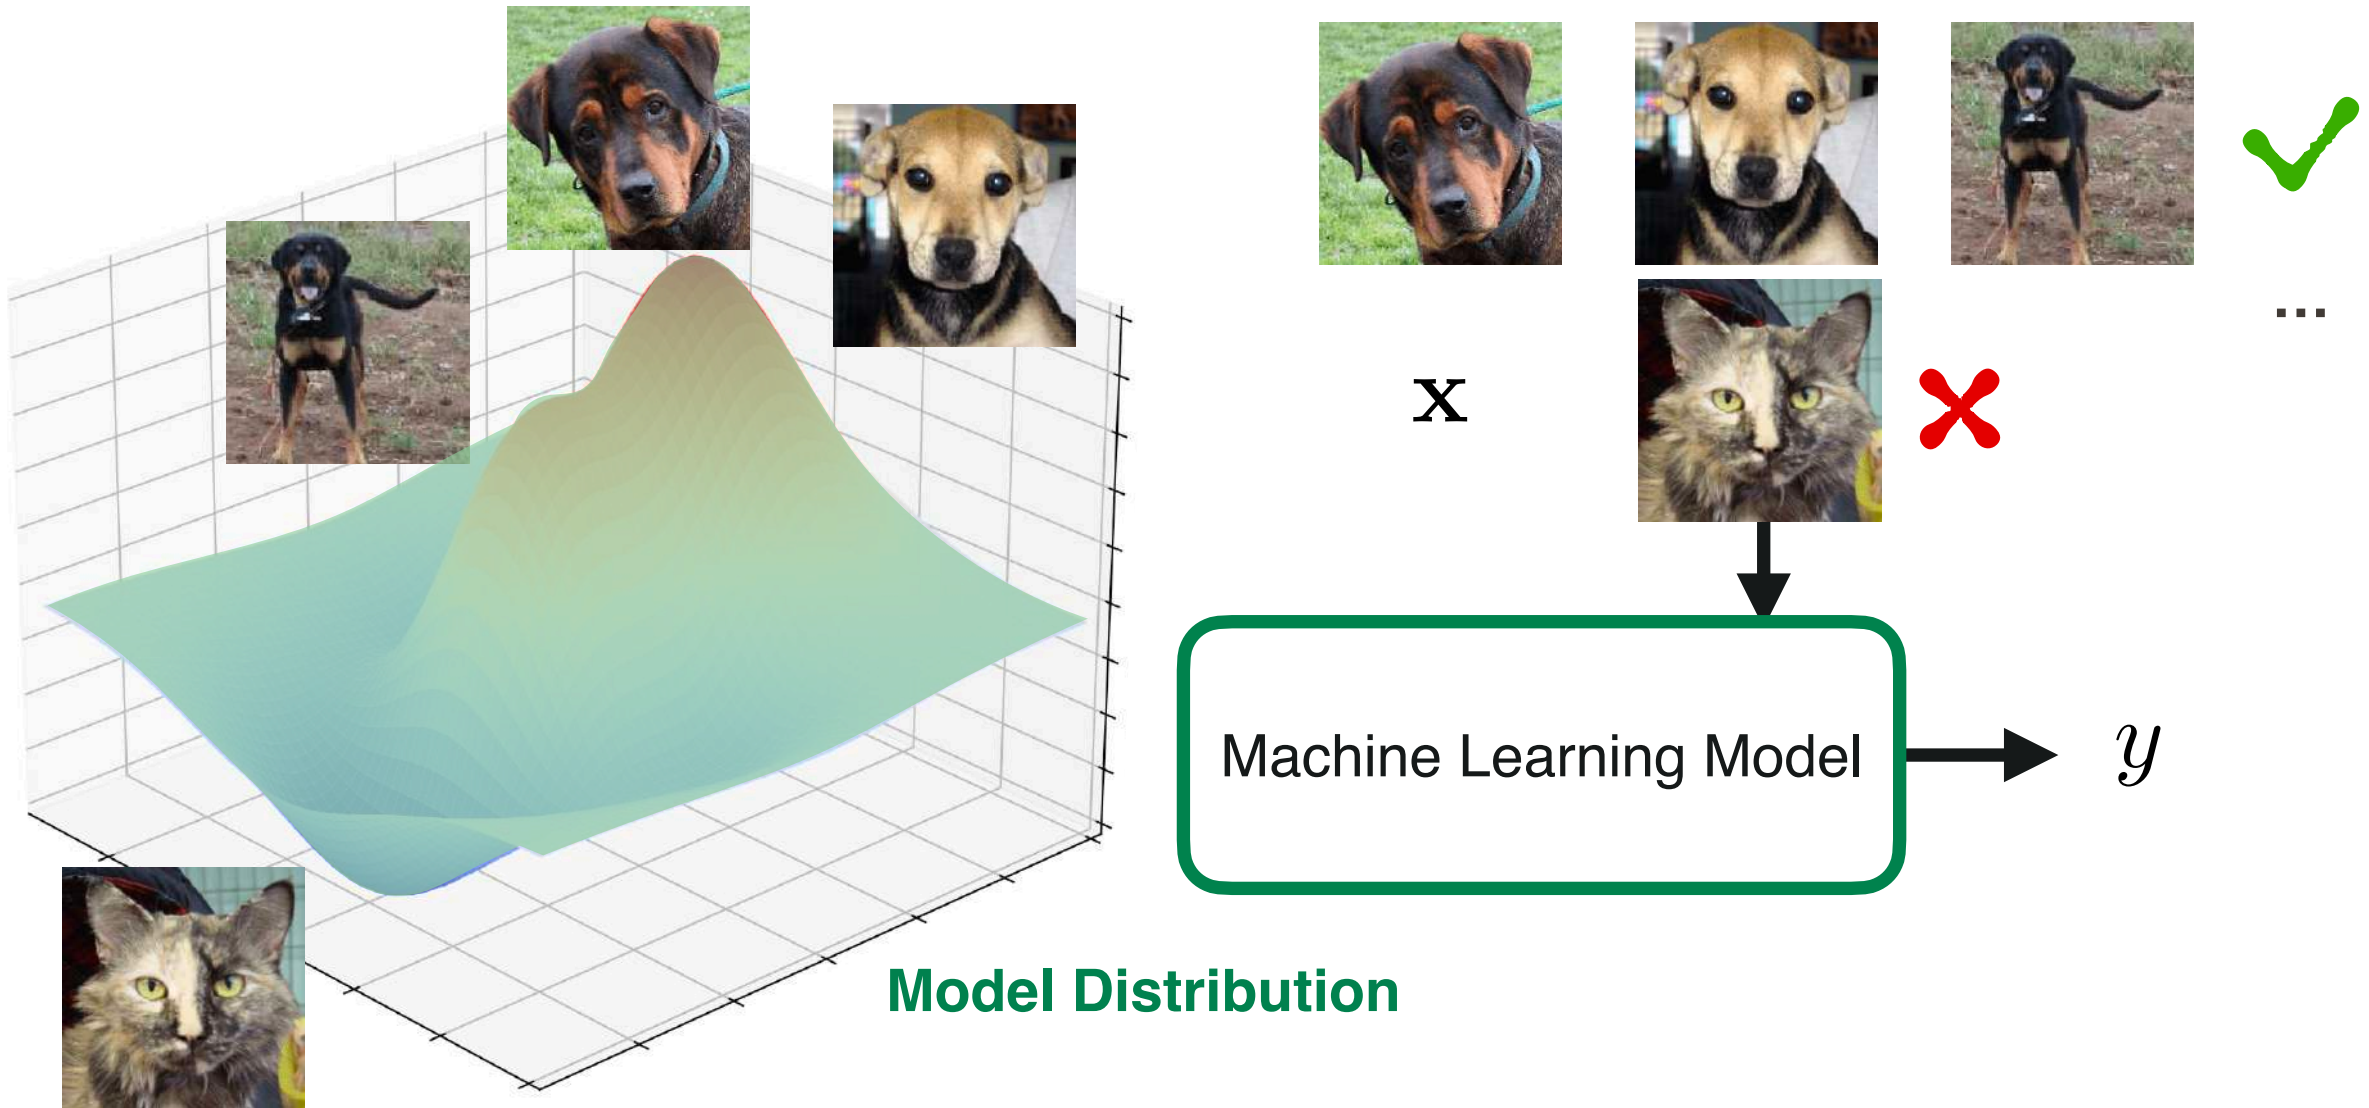


X

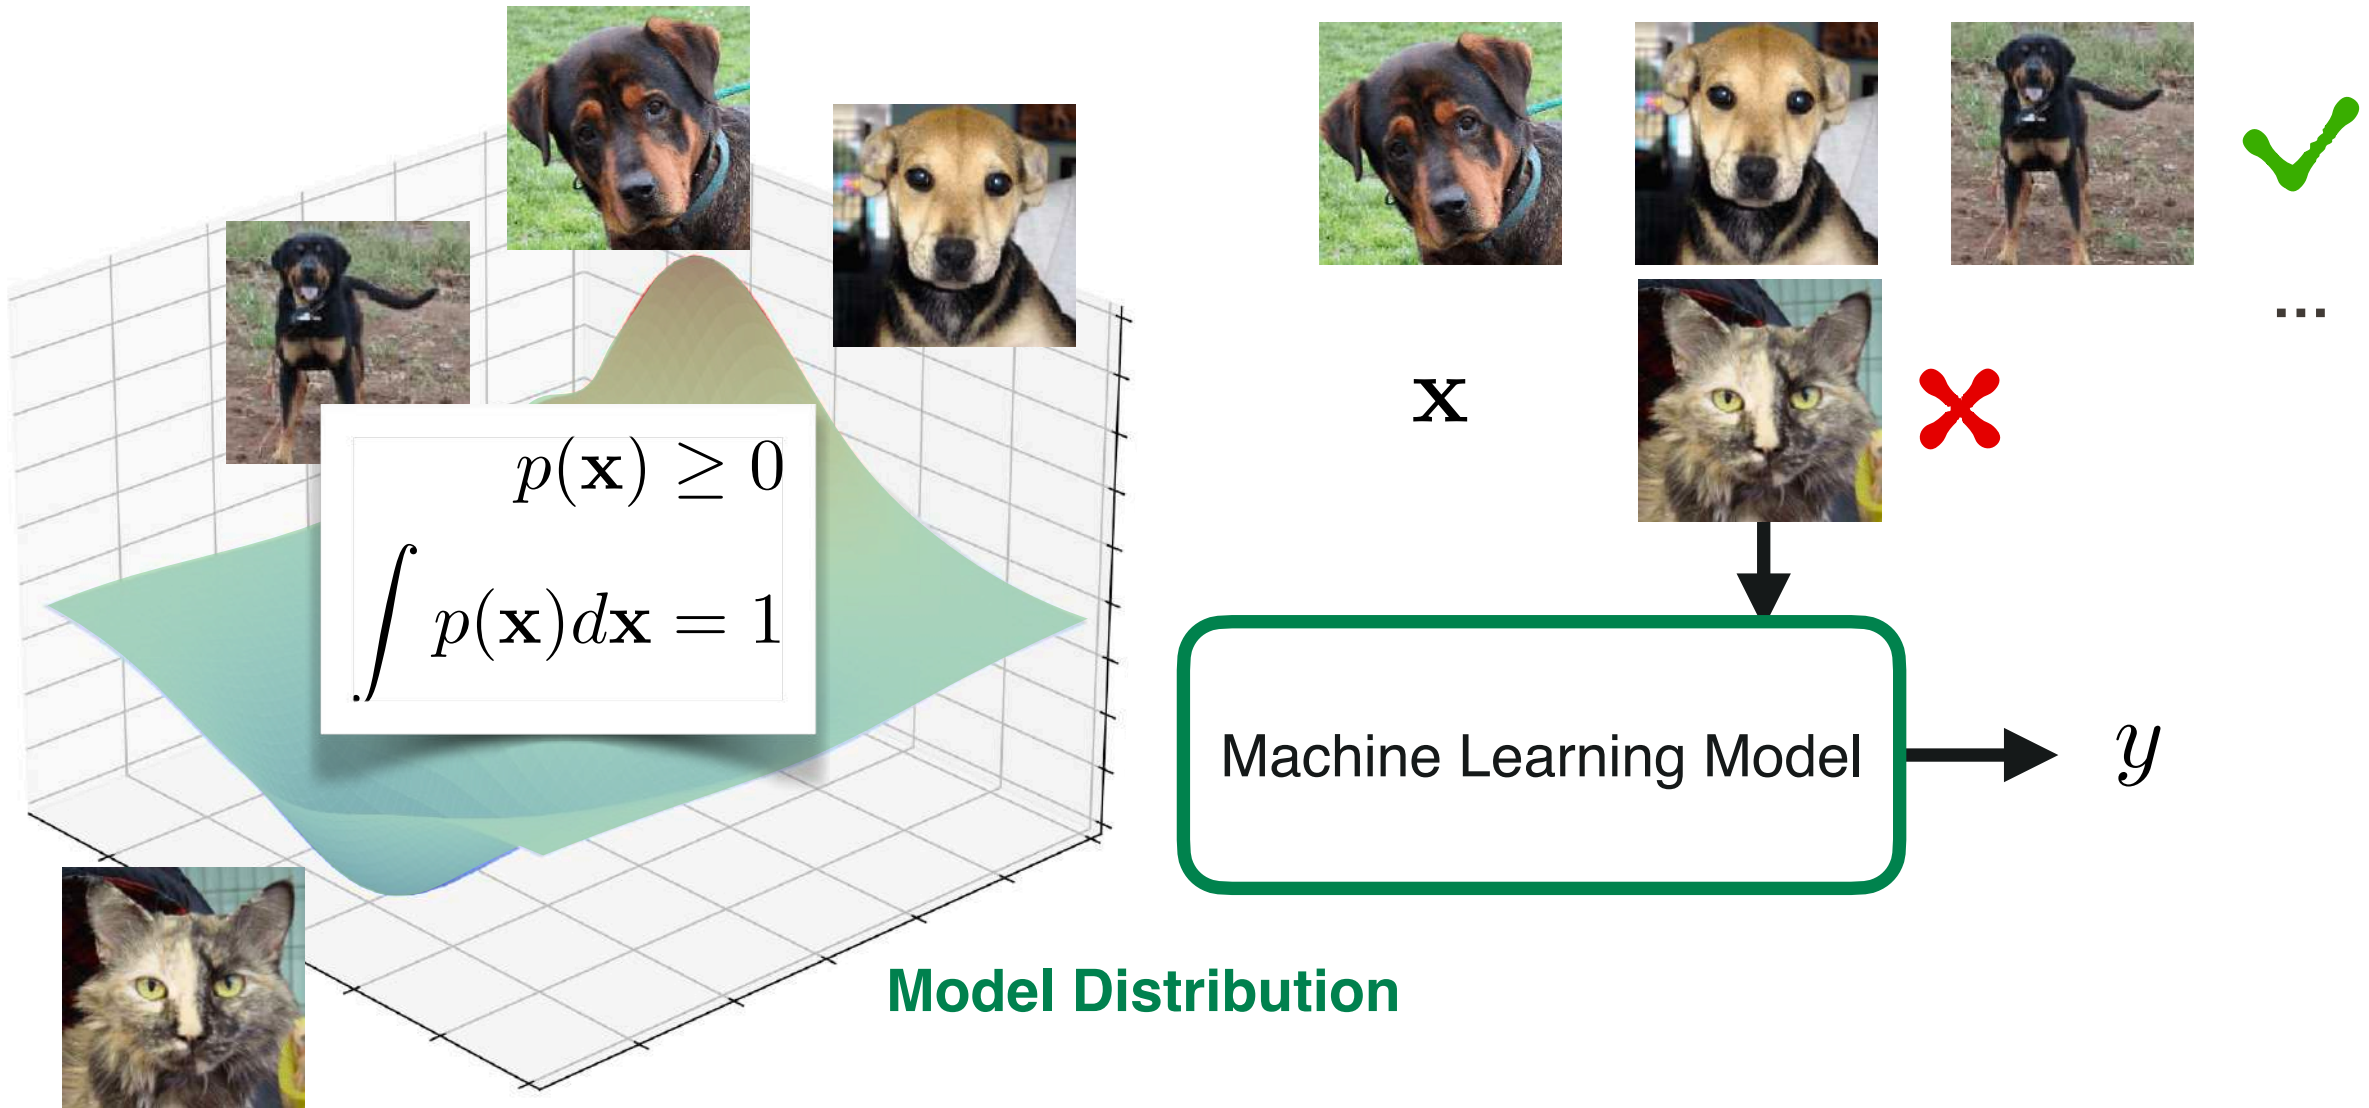


...


Fitting Likelihood-Based Models



Fitting Likelihood-Based Models



Fitting Likelihood-Based Models



Model Distribution



...

\mathbf{x}

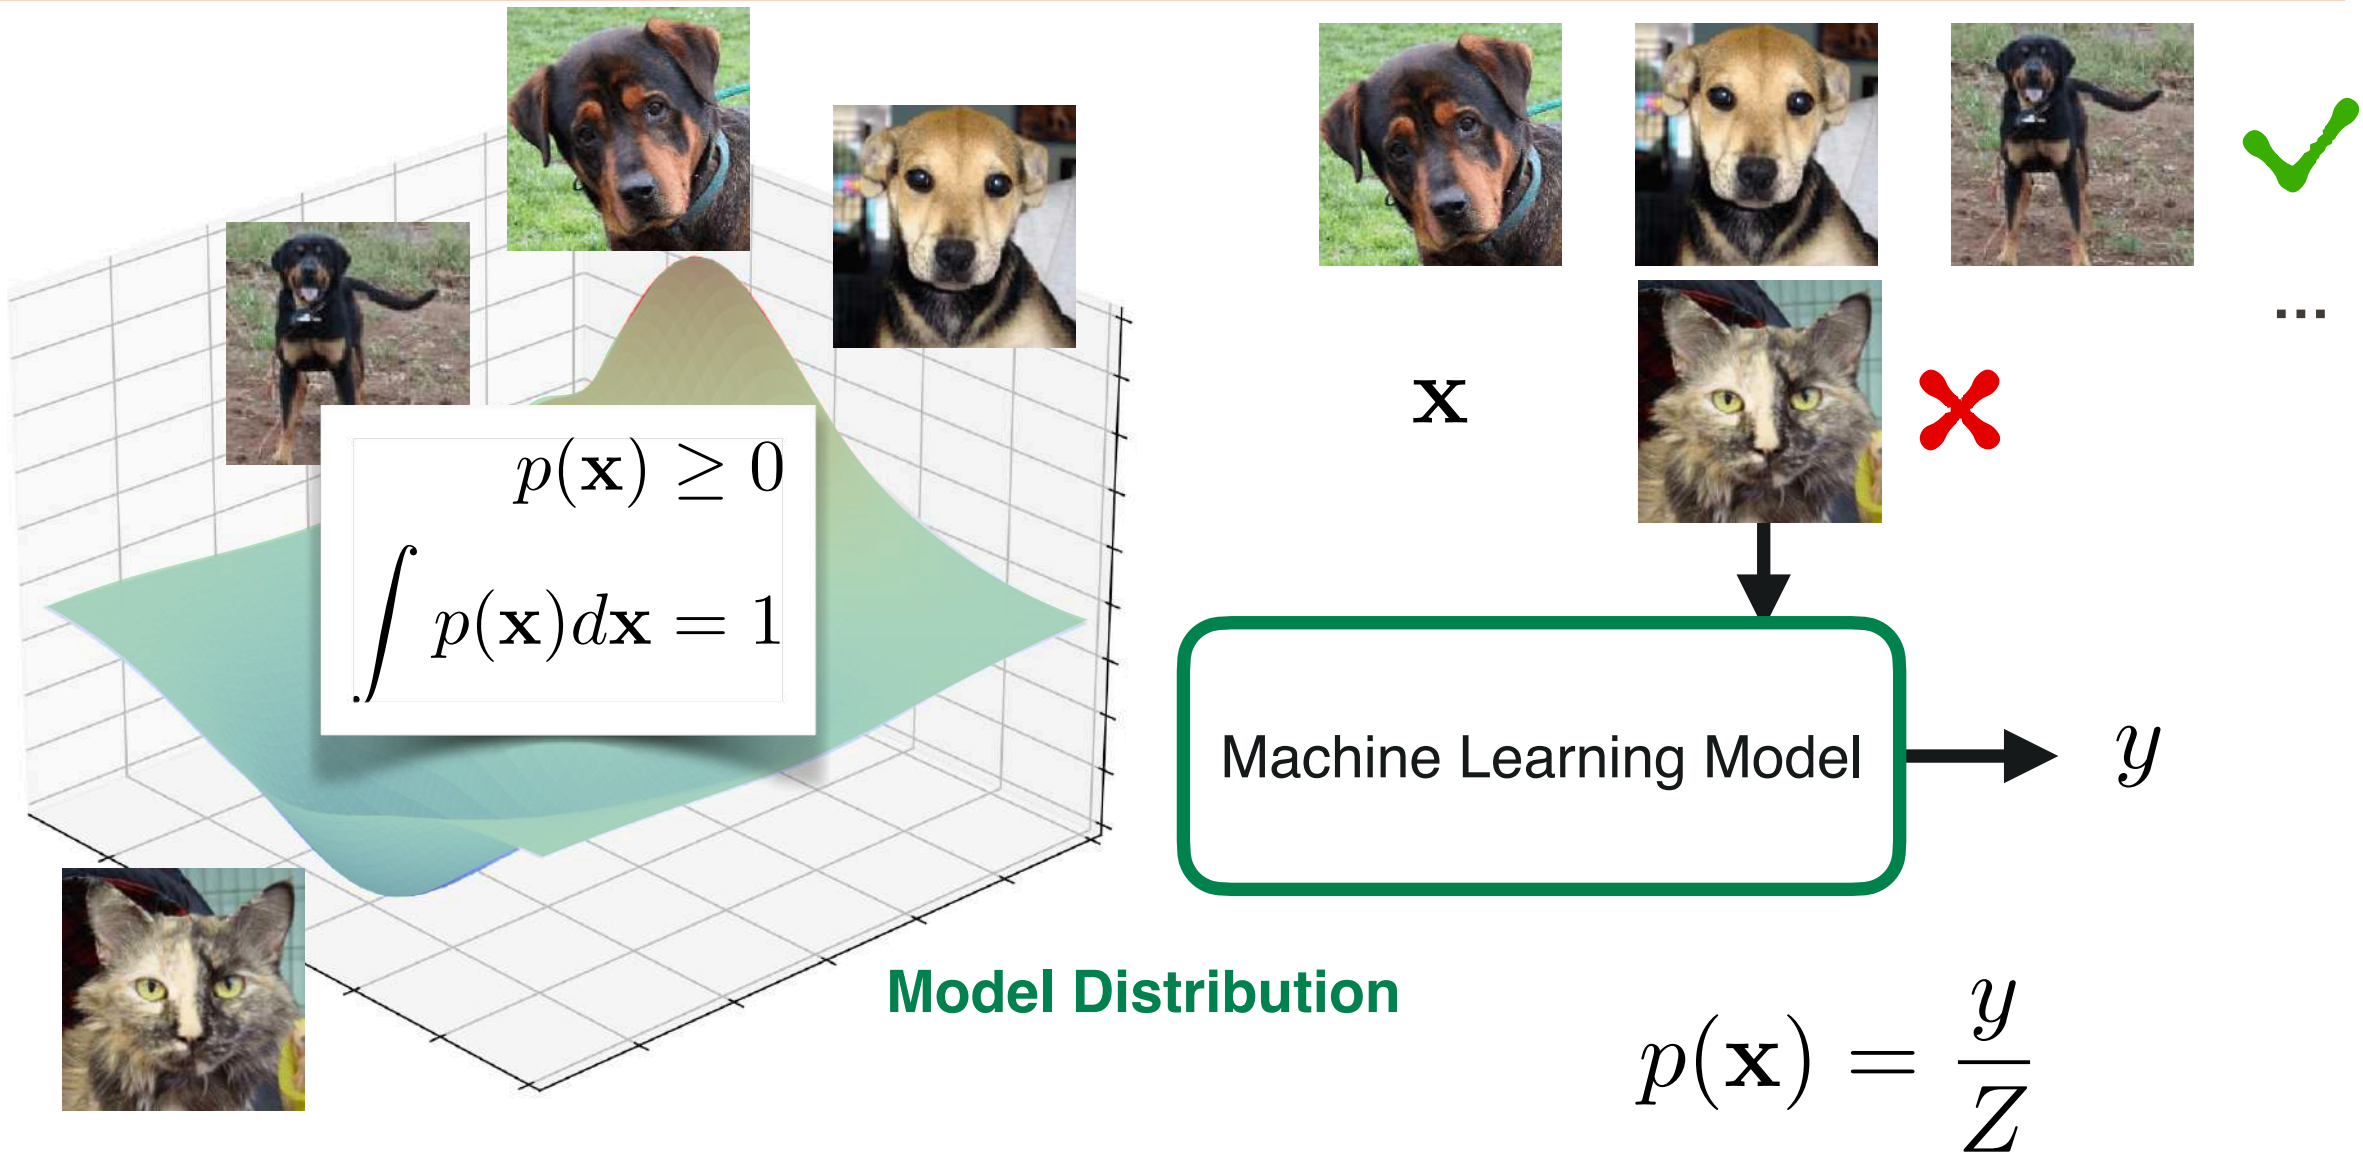


Machine Learning Model

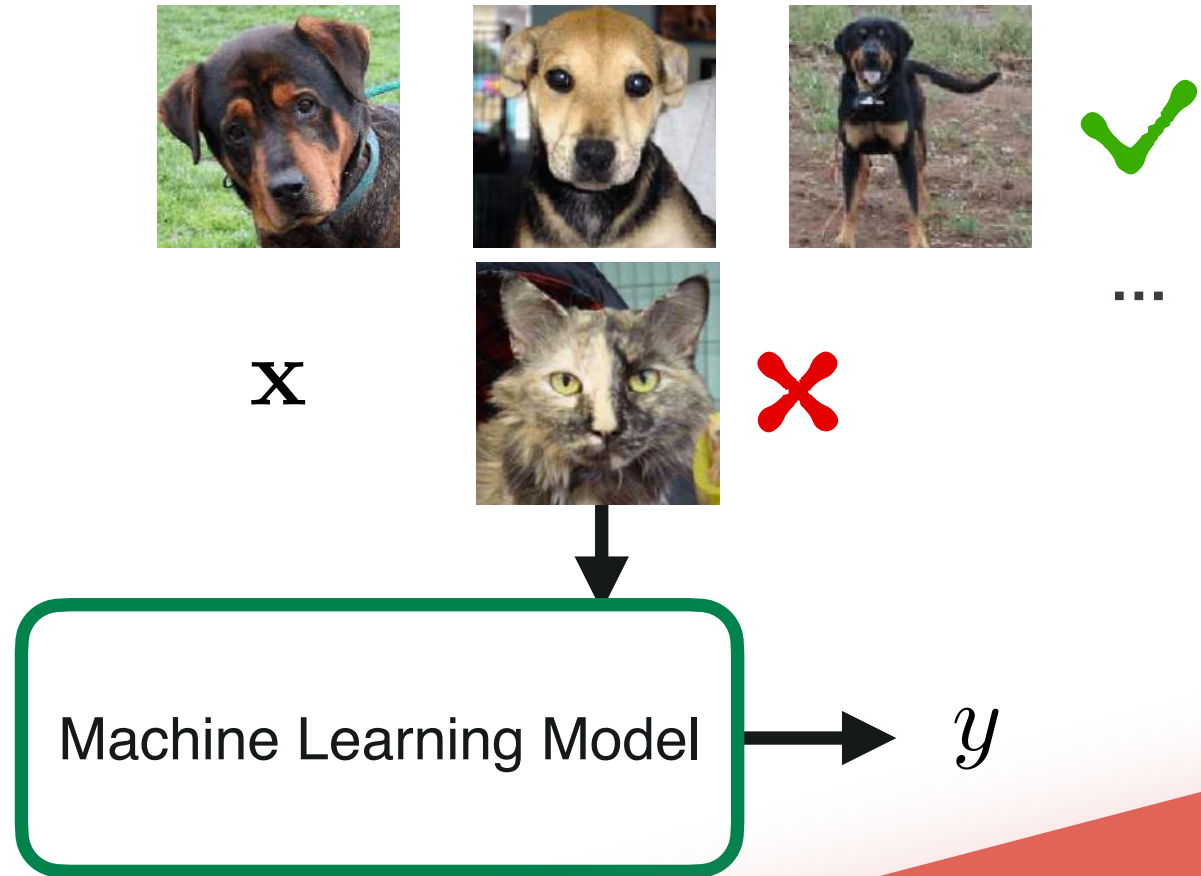

y

$$Z = \int y d\mathbf{x}$$

Fitting Likelihood-Based Models



Fitting Likelihood-Based Models



Intractable!

Fitting Likelihood-Based Models

$$p(\mathbf{x}) = \frac{y}{Z}$$

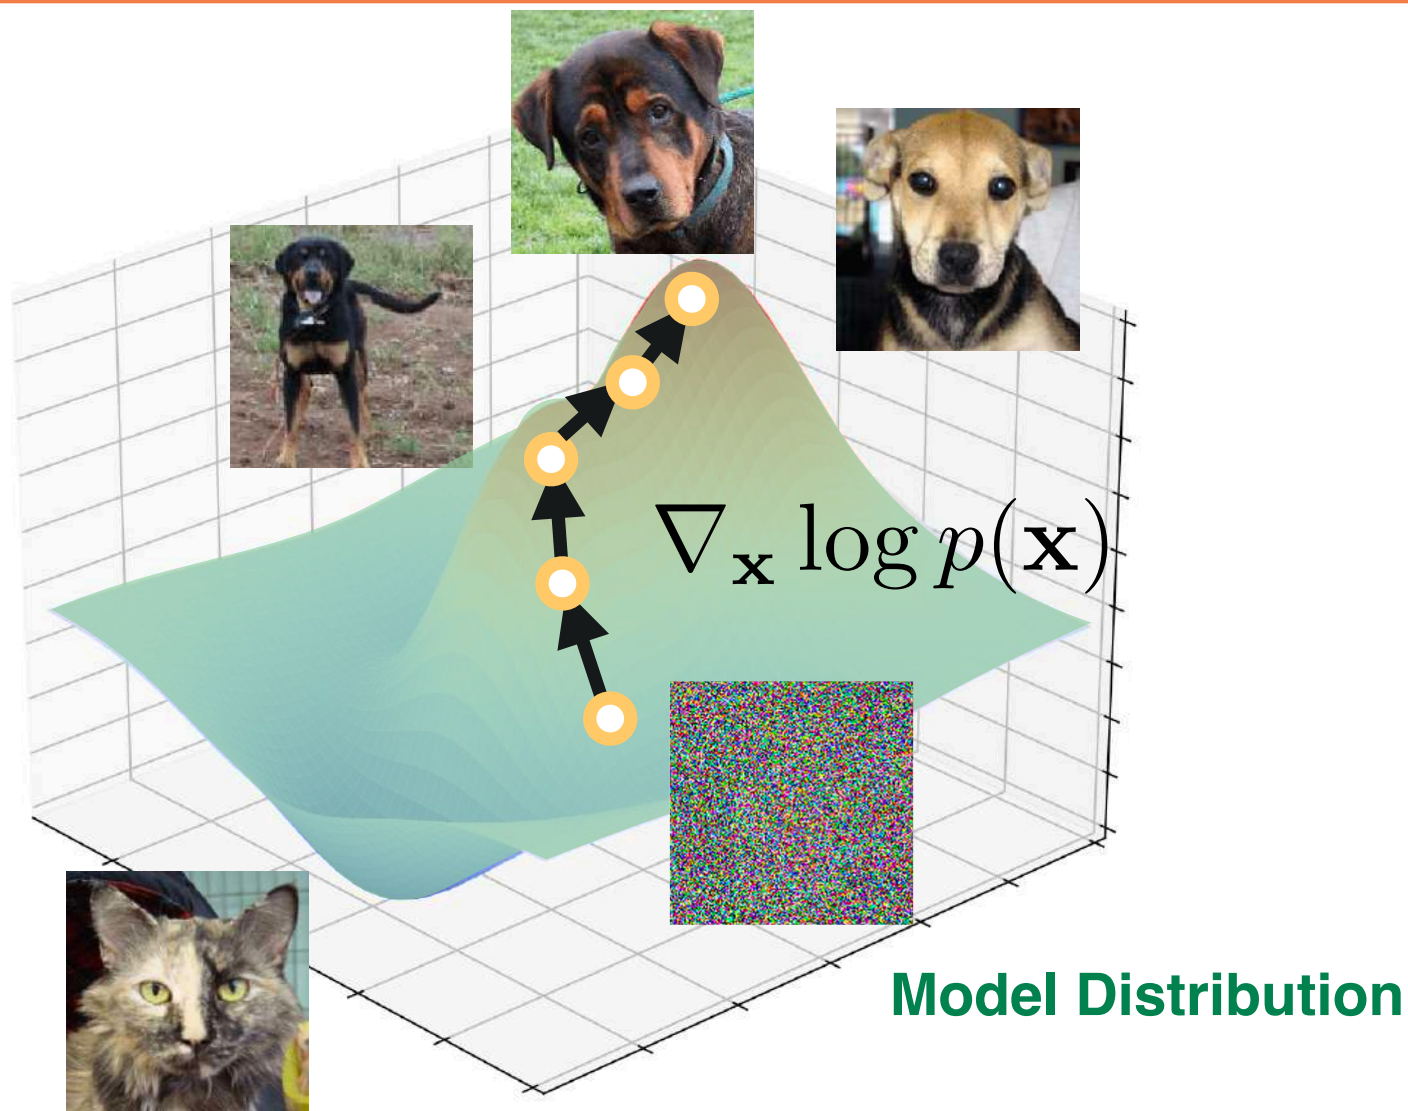
$$\log p(\mathbf{x}) = \log y - \log Z$$

$$\nabla_{\mathbf{x}} \log p(\mathbf{x}) = \nabla_{\mathbf{x}} \log y - \cancel{\nabla_{\mathbf{x}} \log Z} \quad \mathbf{0!}$$

Fit this instead

Machine Learning Model


Fitting Likelihood-Based Models



Why is this called diffusion model?

Forward Diffusion Process

1. Sample a random noise image $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$




2. Add noise by blending. *[This is a designed procedure/a schedule]*


Reverse Diffusion Process

1. Sample a random noise image $\mathbf{x}^T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

How do we get this clean image?



\mathbf{x}^T




\mathbf{x}^0

Reverse Diffusion Process


1. Sample a random noise image $\mathbf{x}^T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

How do we get this clean image?



\mathbf{x}^T

$\epsilon_{\theta}(\mathbf{x}^T, T)$



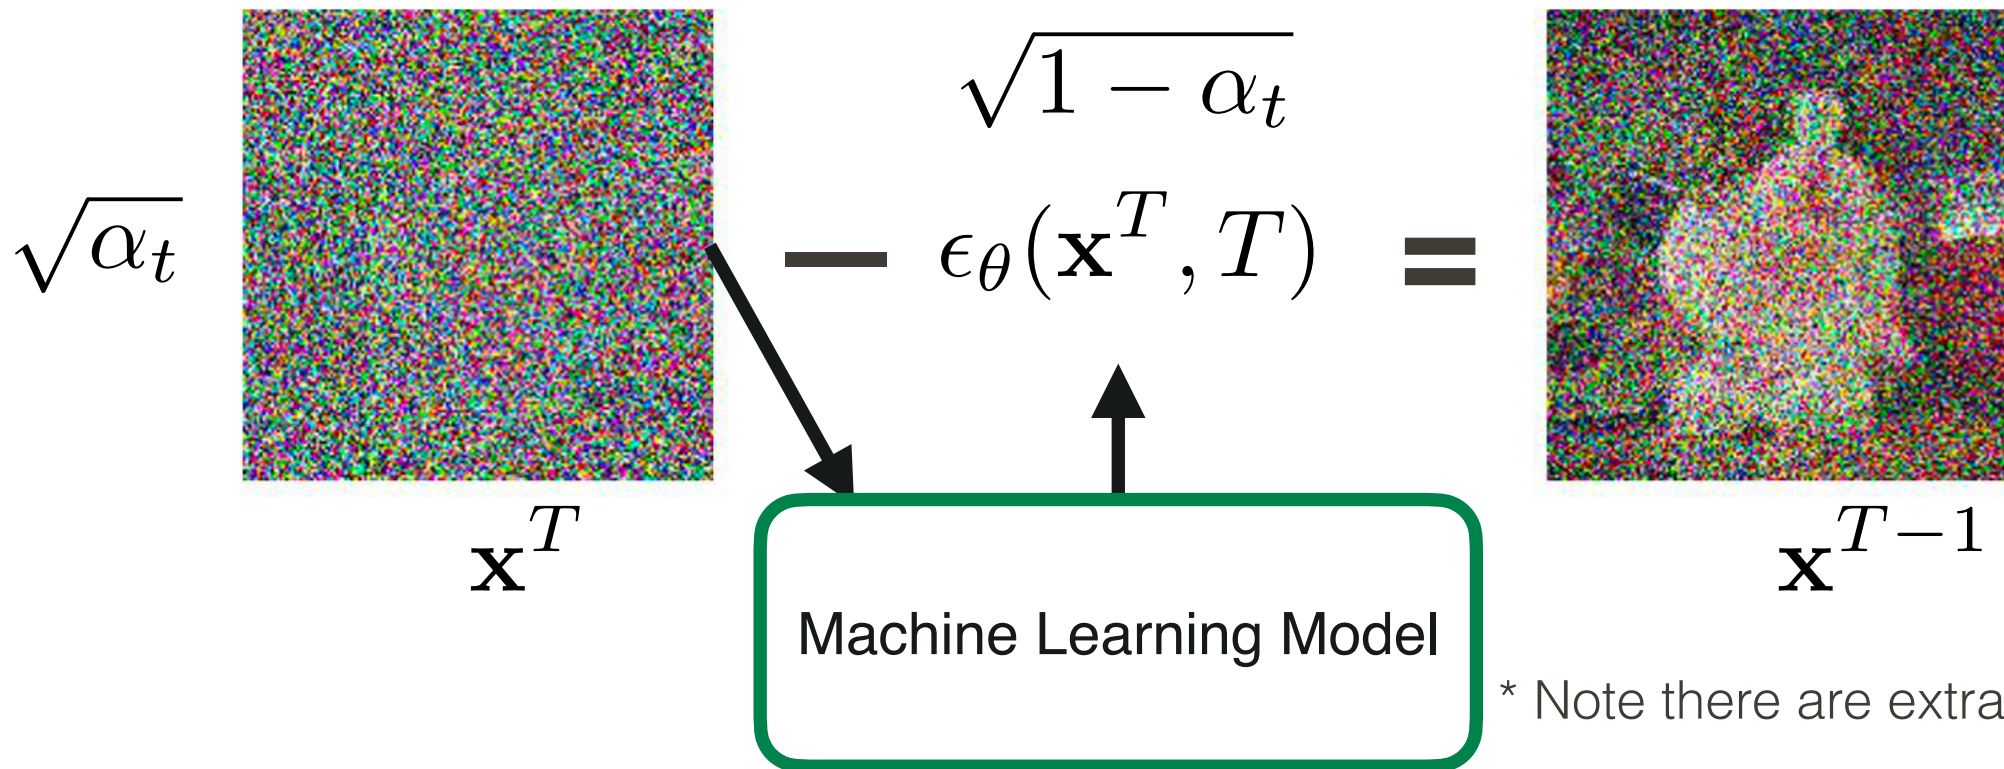
\mathbf{x}^0

Machine Learning Model

Reverse Diffusion Process

1. Sample a random noise image $\mathbf{x}^T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$


How do we get this clean image?




Reverse Diffusion Process

1. Sample a random noise image $\mathbf{x}^T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

2. Denoise



\mathbf{x}^T



\mathbf{x}^{T-1}

Reverse Diffusion Process

1. Sample a random noise image $\mathbf{x}^T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

2. Denoise



\mathbf{x}^{T-1}



\mathbf{x}^{T-2}


Reverse Diffusion Process

1. Sample a random noise image $\mathbf{x}^T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

2. Denoise




\mathbf{x}^{T-2}




\mathbf{x}^0

Model Fitting

Machine Learning Model?



\mathbf{x}^0



\mathbf{x}^T

ϵ

$$\mathbb{E}_{t, \mathbf{x}^0, \epsilon_t} \|\epsilon_\theta(\mathbf{x}^t, t) - \epsilon_t\|^2$$

Connection between Two Definitions

Want to fit $\nabla_{\mathbf{x}} \log p(\mathbf{x})$
with $s_{\theta}(\mathbf{x})$

Connection between Two Definitions



Want to fit $\nabla_{\mathbf{x}} \log p(\mathbf{x})$
with $s_{\theta}(\mathbf{x})$

$p(\mathbf{x})$ is unknown!

Connection between Two Definitions

Want to fit $\nabla_{\mathbf{x}} \log p(\mathbf{x})$
with $s_{\theta}(\mathbf{x})$



$p(\mathbf{x})$ is unknown!



Connection between Two Definitions

Want to fit $\nabla_{\mathbf{x}} \log p(\mathbf{x})$
with $s_{\theta}(\mathbf{x})$

$p(\mathbf{x})$ is unknown!





$$\mathbb{E}_{\tilde{\mathbf{x}}, \mathbf{x}} ||s_{\theta}(\tilde{\mathbf{x}}) - \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})||^2$$

Connection between Two Definitions

Want to fit $\nabla_{\mathbf{x}} \log p(\mathbf{x})$
with $s_{\theta}(\mathbf{x})$

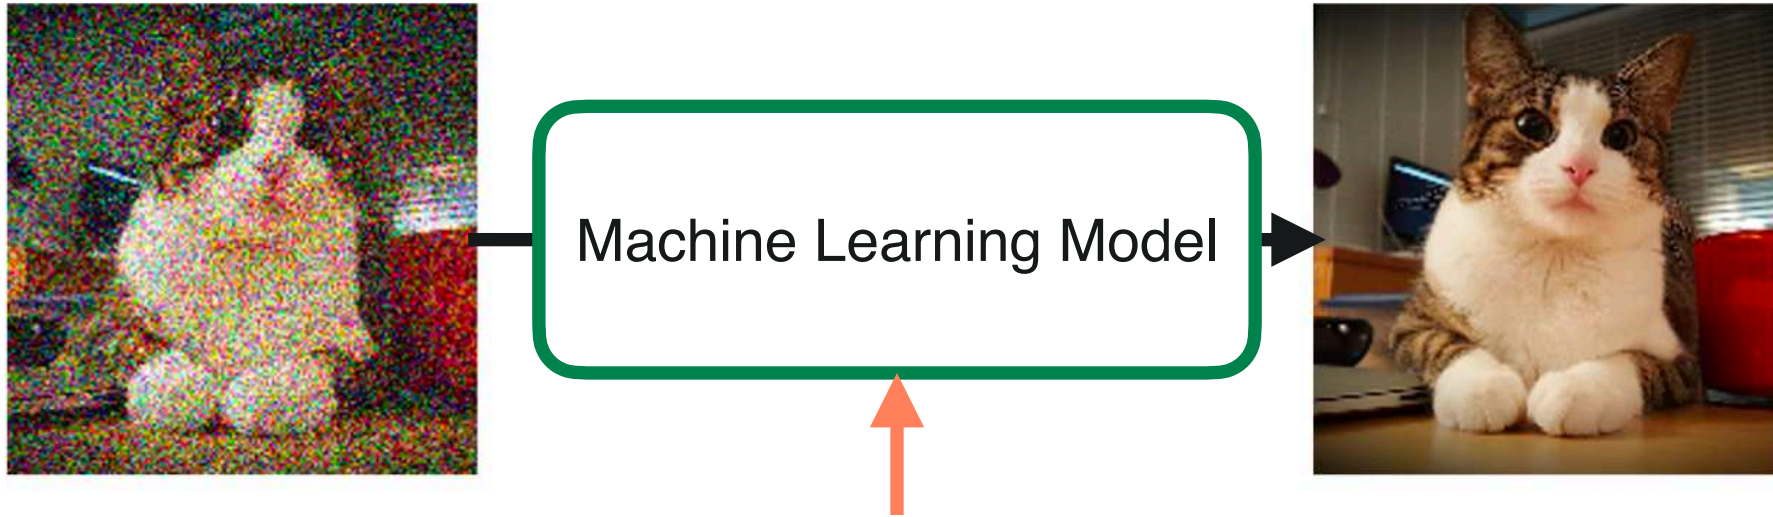
$p(\mathbf{x})$ is unknown!



$$\mathbb{E}_{\tilde{\mathbf{x}}, \mathbf{x}} ||s_{\theta}(\tilde{\mathbf{x}}) - \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})||^2$$

$$\mathbb{E}_{t, \mathbf{x}^0, \epsilon_t} ||\epsilon_{\theta}(\mathbf{x}^t, t) - \epsilon_t||^2$$


Conditional v.s Unconditional



"A photo of a cat perching on a desk"

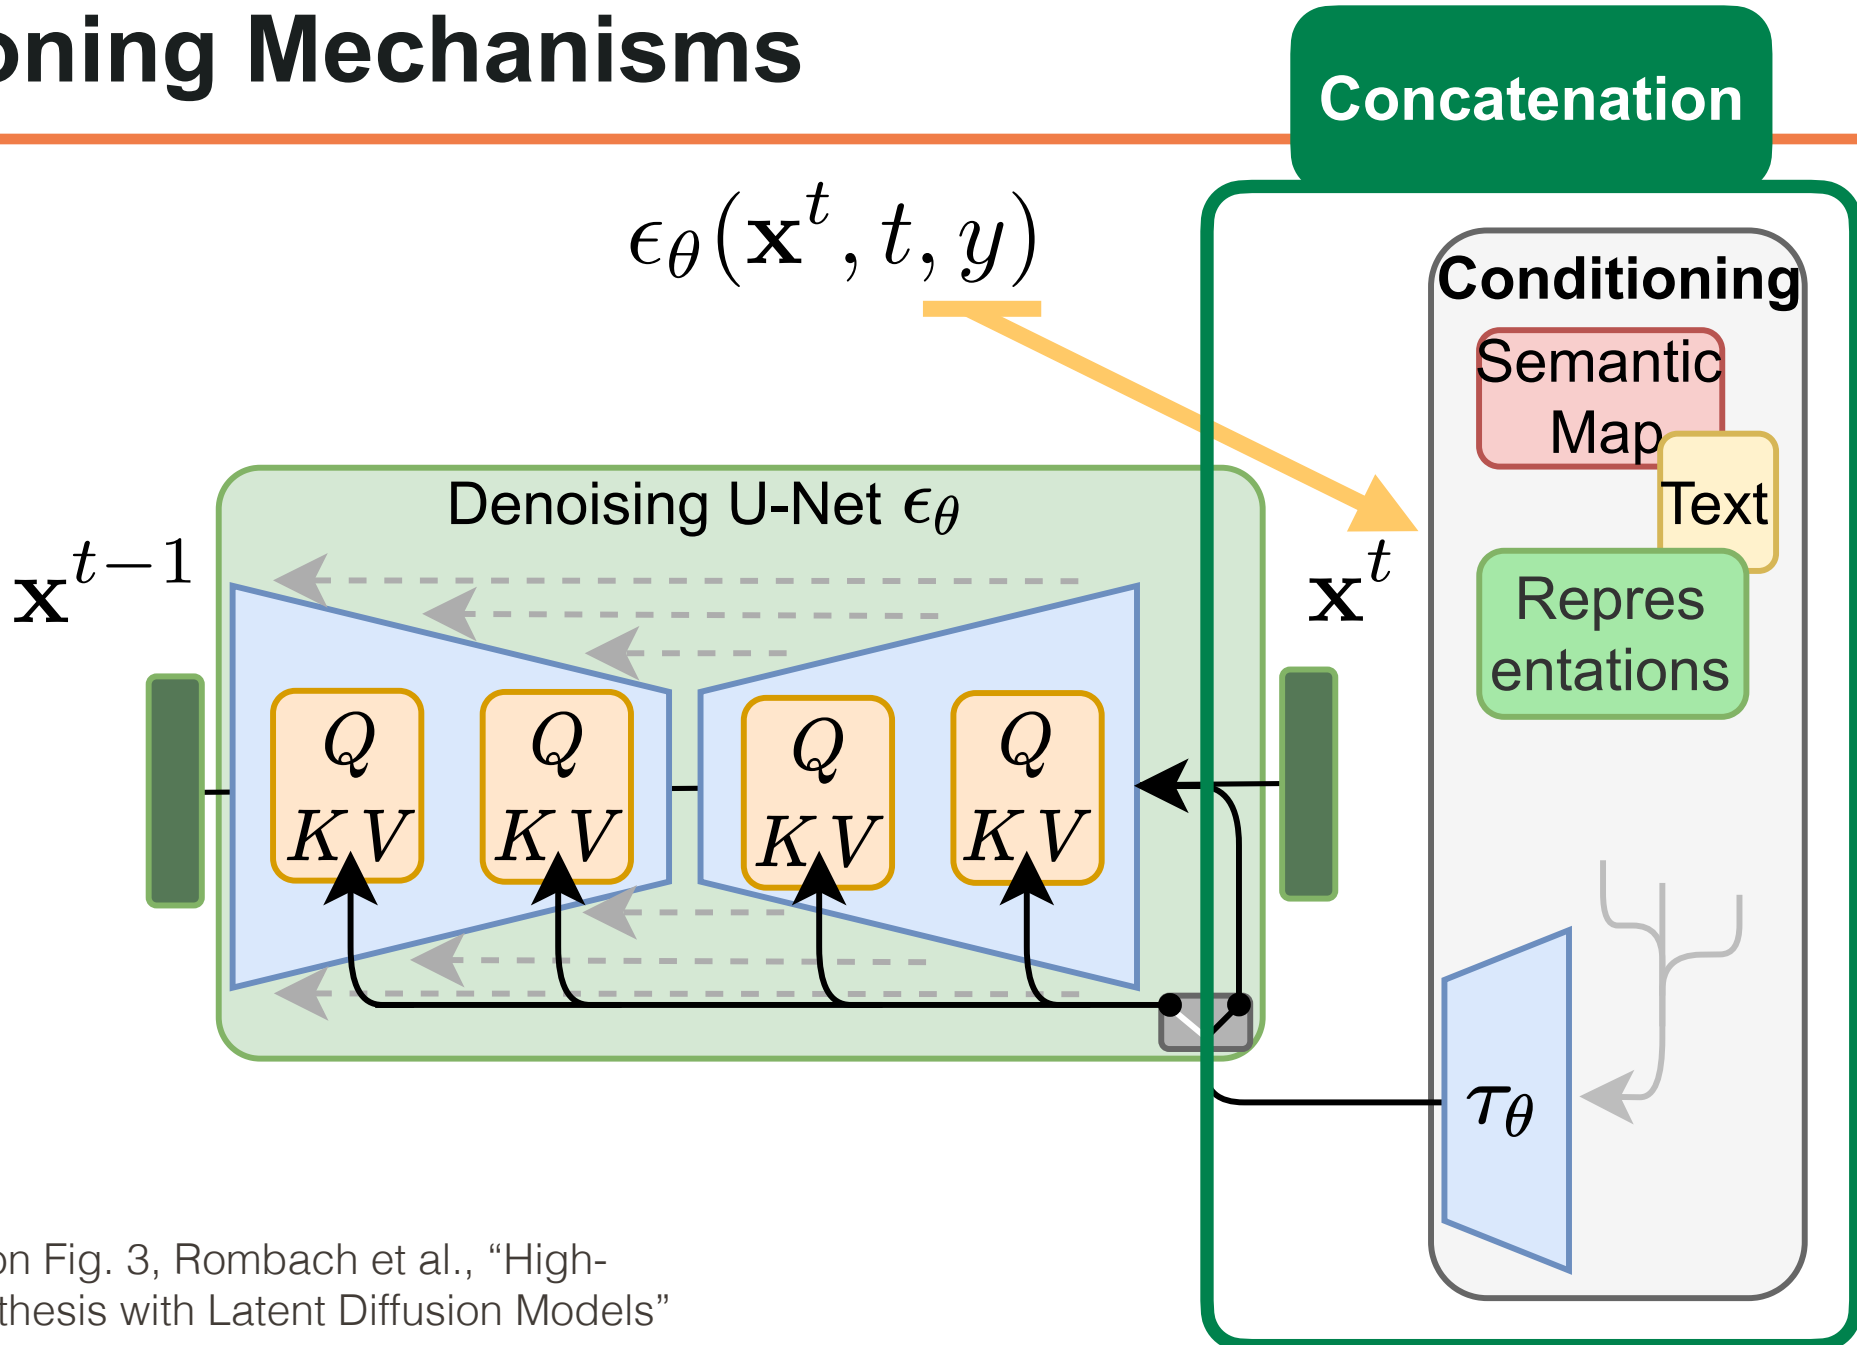
- More control for users.
In contrast, unconditional generative model would need to use random seeds to control the output.
- Empirically, conditional generative models are **easier to train and perform better** than unconditional ones.

Conditioning Mechanisms




* Modified based on Fig. 3, Rombach et al., “High-Resolution Image Synthesis with Latent Diffusion Models”

Conditioning Mechanisms



* Modified based on Fig. 3, Rombach et al., "High-Resolution Image Synthesis with Latent Diffusion Models"

Conditioning Mechanisms



* Modified based on Fig. 3, Rombach et al., "High-Resolution Image Synthesis with Latent Diffusion Models"

Cross-Attention

Q: Query

“ \mathbf{x}^t ”

K: Key

“ y ”

V: Value

“ y ”

Cross-Attention

Q: Query

“ \mathbf{x}^t ”

K: Key

“ y ”

V: Value

“ y ”

$$\text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$

Cross-Attention

Q: Query

“ \mathbf{x}^t ”

K: Key


“ y ”

V: Value

“ y ”

Similarity Score

$$\text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$



Cross-Attention

Q: Query

“ \mathbf{x}^t ”

K: Key

“ y ”

V: Value

“ y ”

Normalization

$$\text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$

Cross-Attention

Q: Query

“ \mathbf{x}^t ”

K: Key

“ y ”

V: Value

“ y ”

$$\text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$

Condition is kept in places where **condition and image is similar.**

Cross-Attention

Q: Query

“ \mathbf{x}^t ”

K: Key

“ y ”

V: Value

“ y ”

$$\text{“}\mathbf{x}^t\text{”} + \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$

Cross-Attention


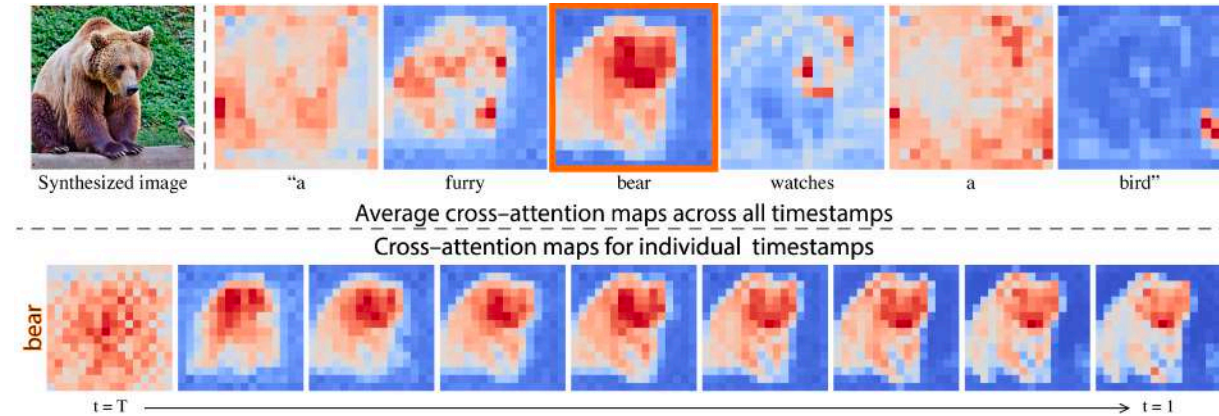


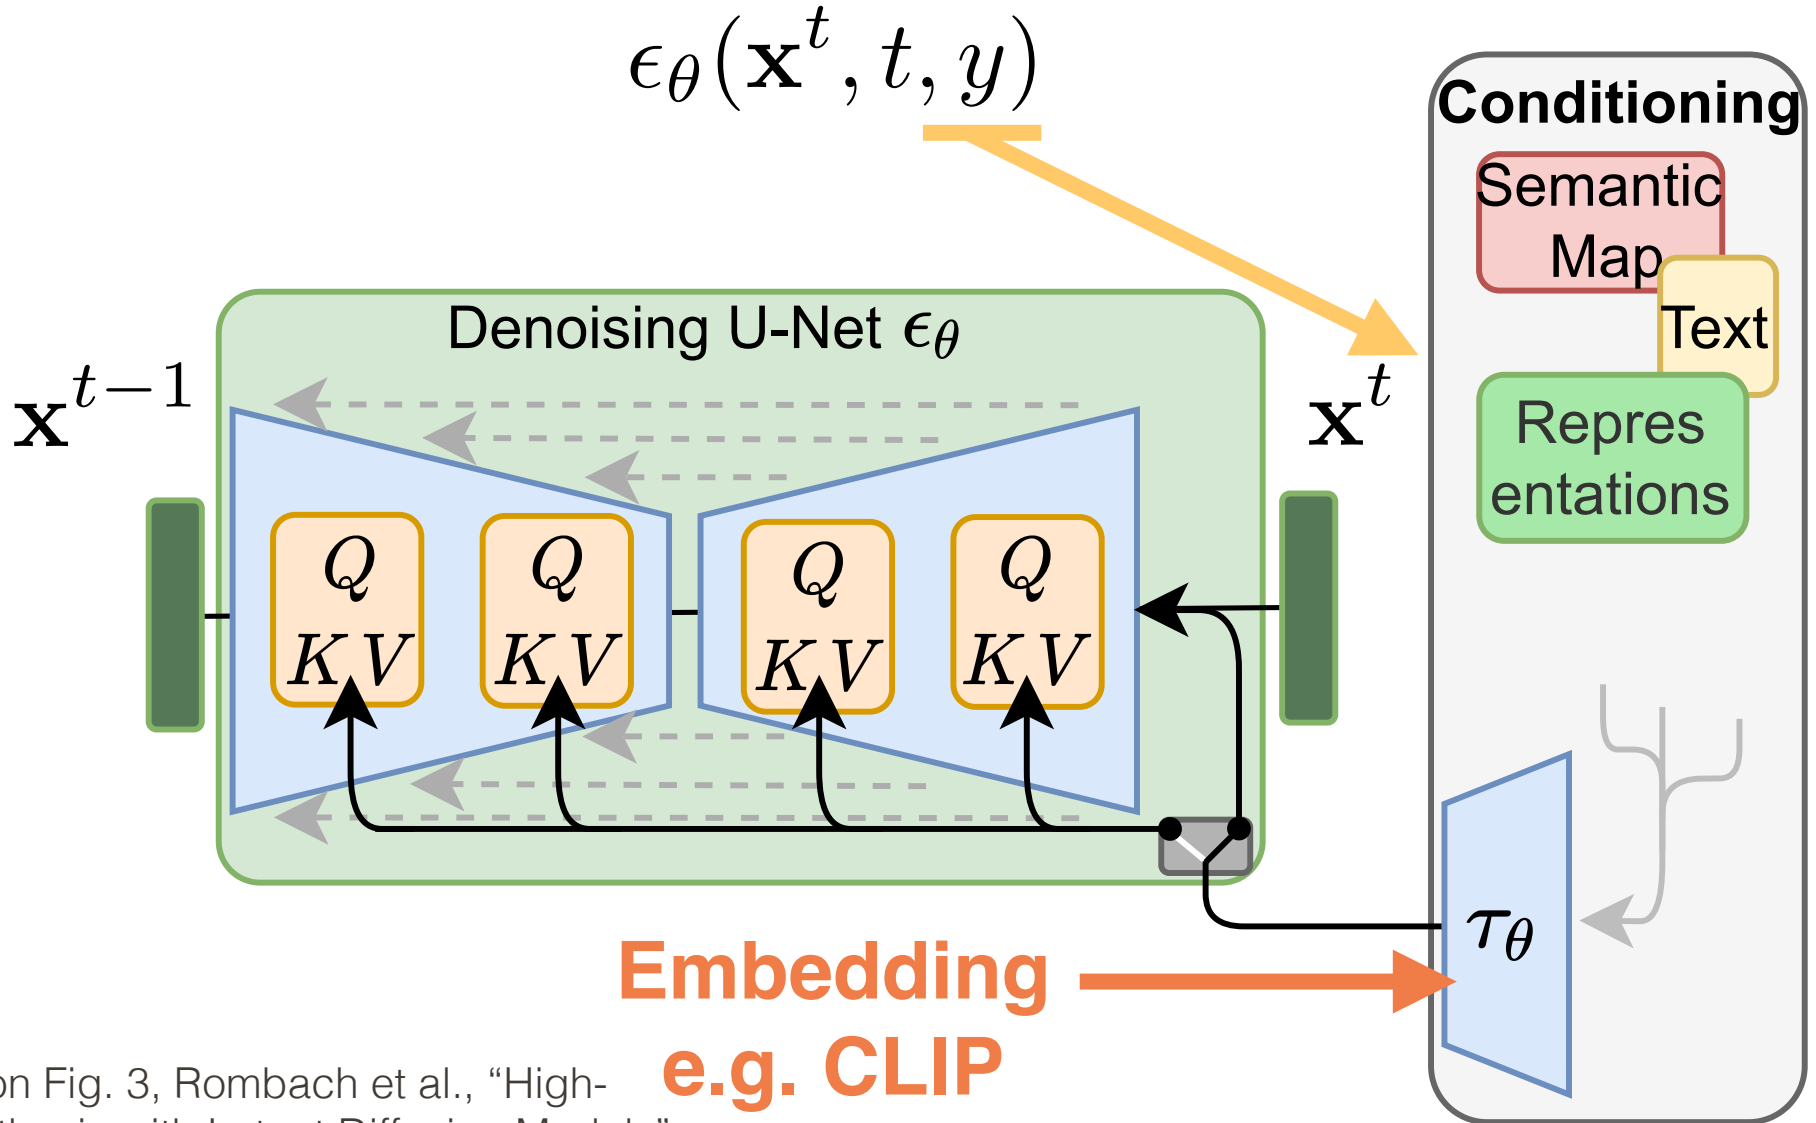
Figure 1: The original synthesized image and three DAAM maps for “monkey,” “hat,” and “walking,” from the prompt, “monkey with hat walking.”

Tang et al., “What the DAAM: Interpreting Stable Diffusion Using Cross Attention”



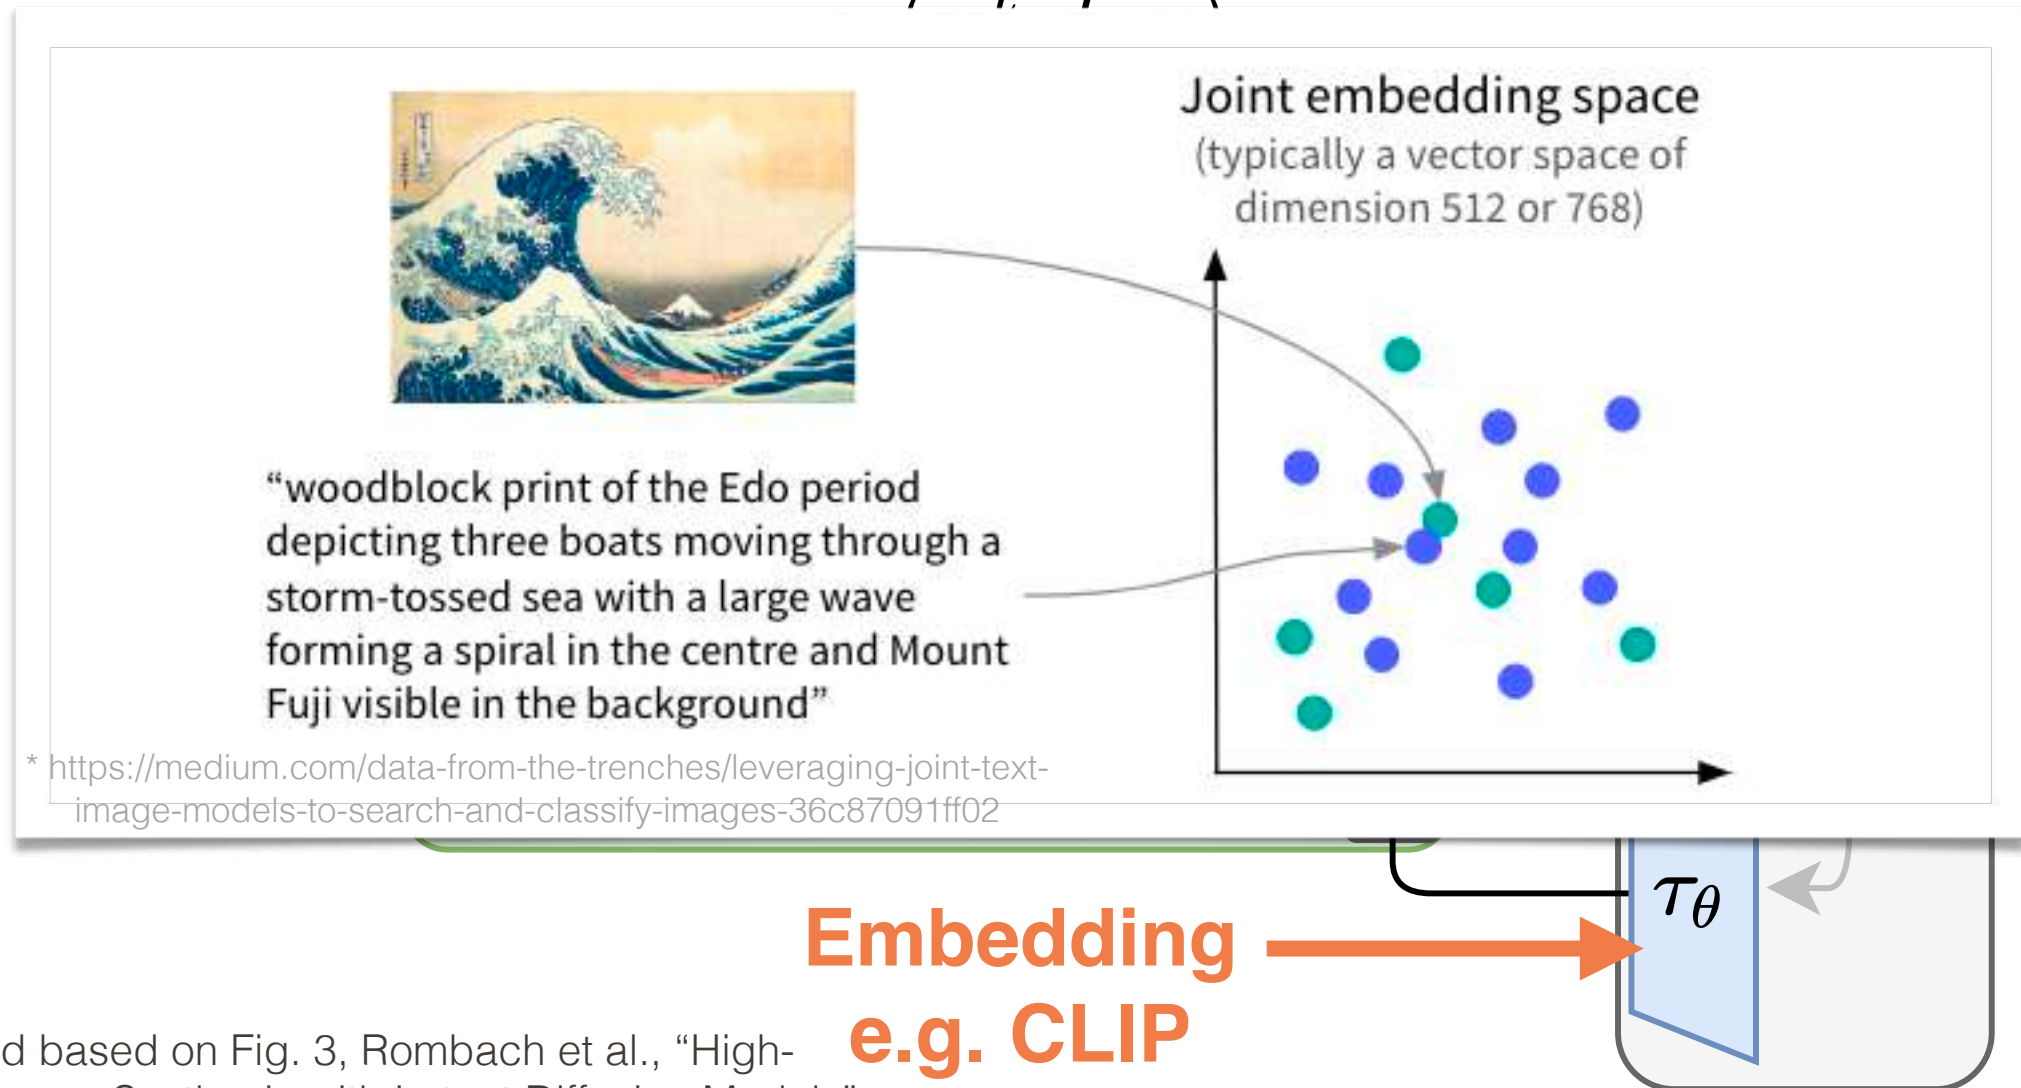
Hertz et al., “Prompt-to-Prompt Image Editing with Cross-Attention Control”

Text Embedding



* Modified based on Fig. 3, Rombach et al., "High-Resolution Image Synthesis with Latent Diffusion Models"

Text Embedding




* Modified based on Fig. 3, Rombach et al., “High-Resolution Image Synthesis with Latent Diffusion Models”

Text Embedding



“woodblock print of the Edo period depicting three boats moving through a storm-tossed sea with a large wave forming a spiral in the centre and Mount Fuji visible in the background”

* <https://medium.com/data-from-the-trenches/leveraging-joint-text-image-models-to-search-and-classify-images-36c87091ff02>



Tyshchuk et al., “On Isotropy of Multimodal Embeddings”

Embedding
e.g. CLIP

τ_θ

* Modified based on Fig. 3, Rombach et al., “High-Resolution Image Synthesis with Latent Diffusion Models”

Questions?

Advanced Diffusion-Model-Based Editing Tools

Image-to-Image Methods

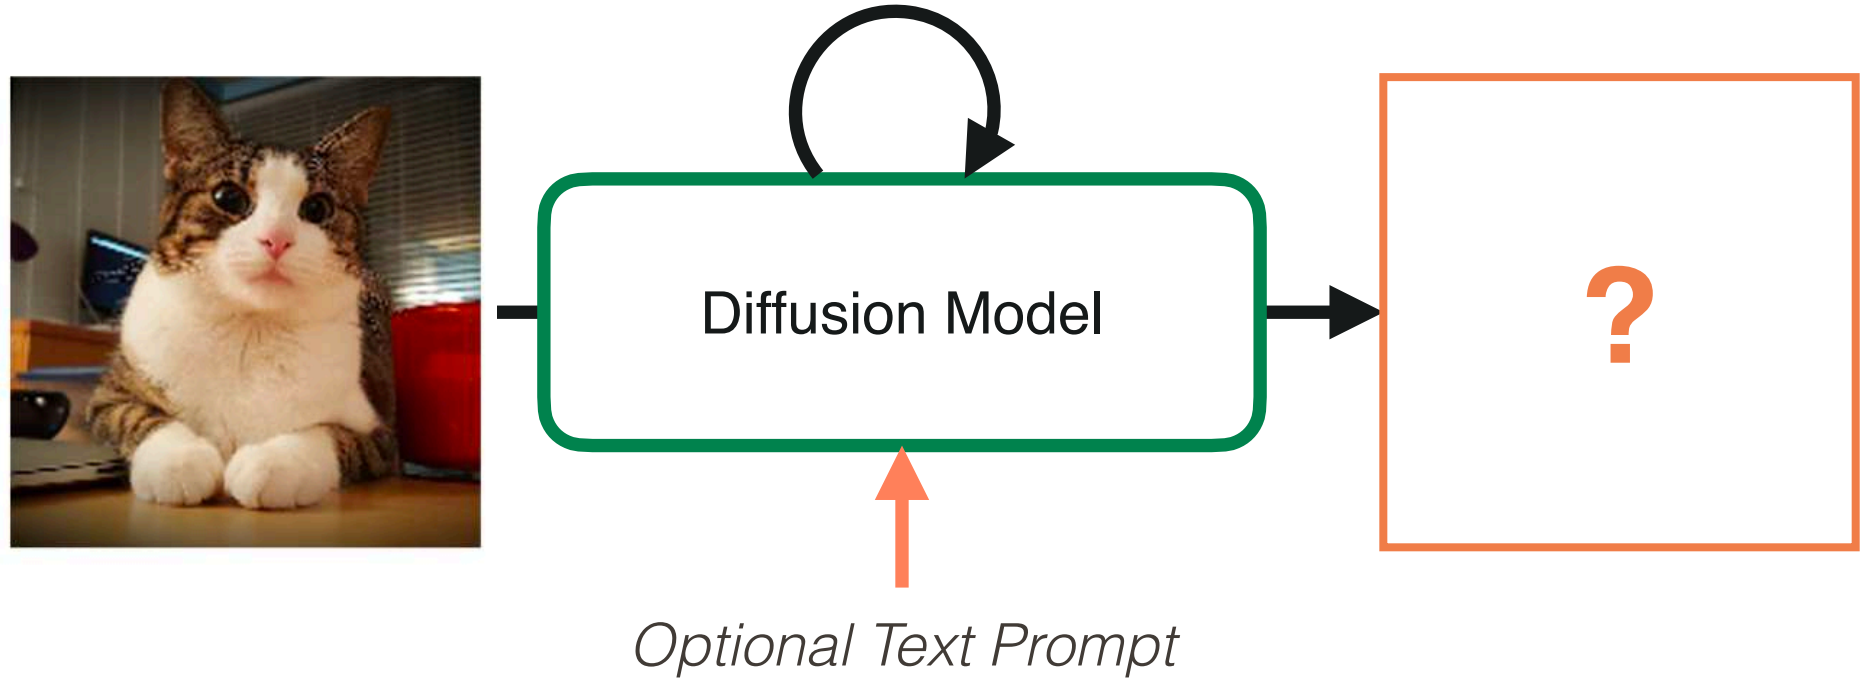


Image-to-Image Methods: Image Variances

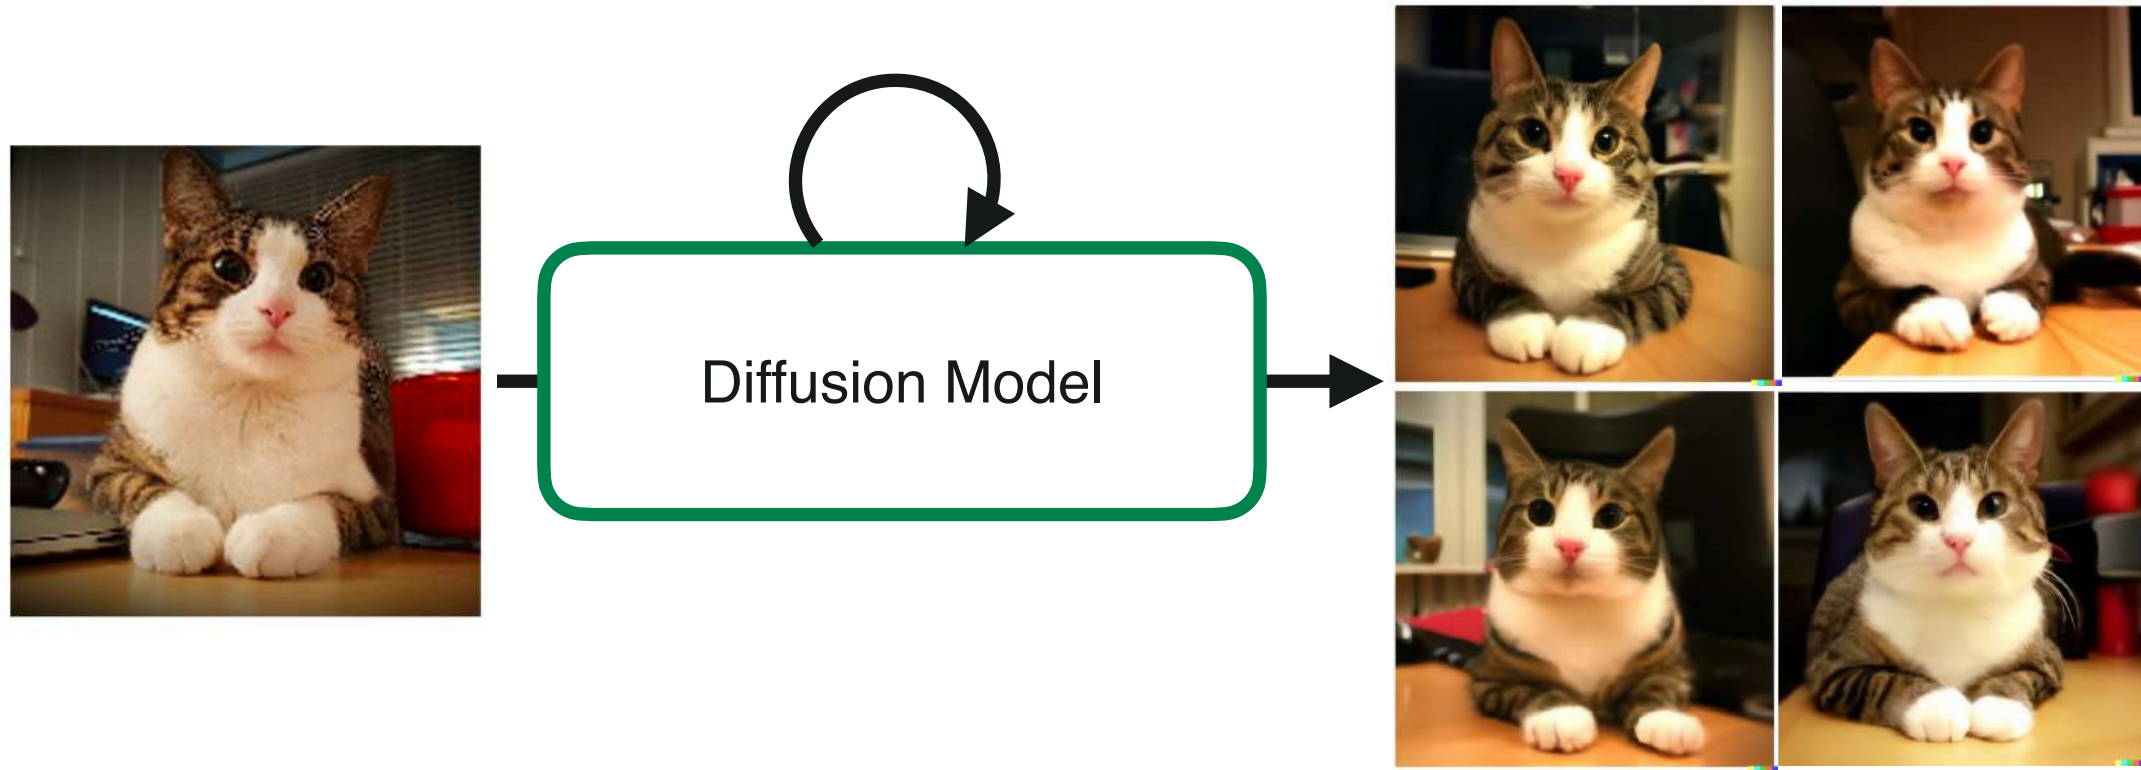


Image-to-Image Methods: Image Variances

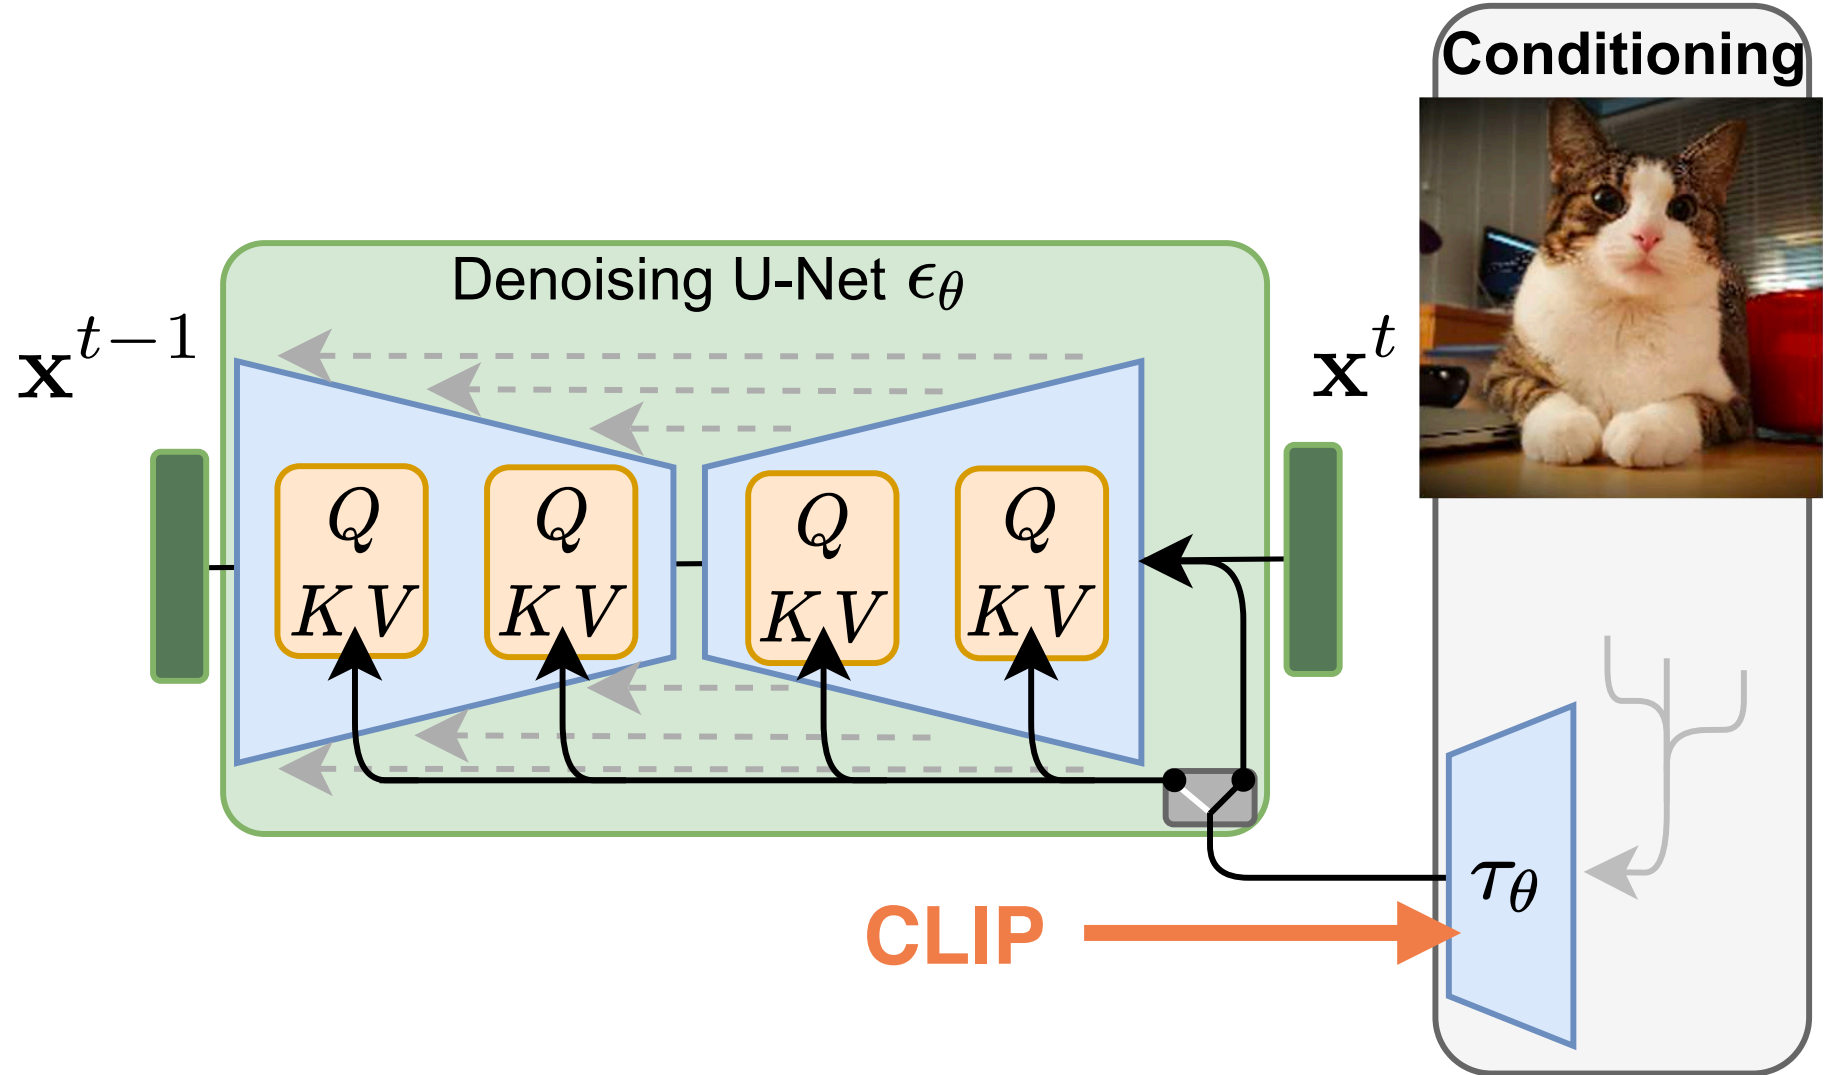


Image-to-Image Methods: ControlNet

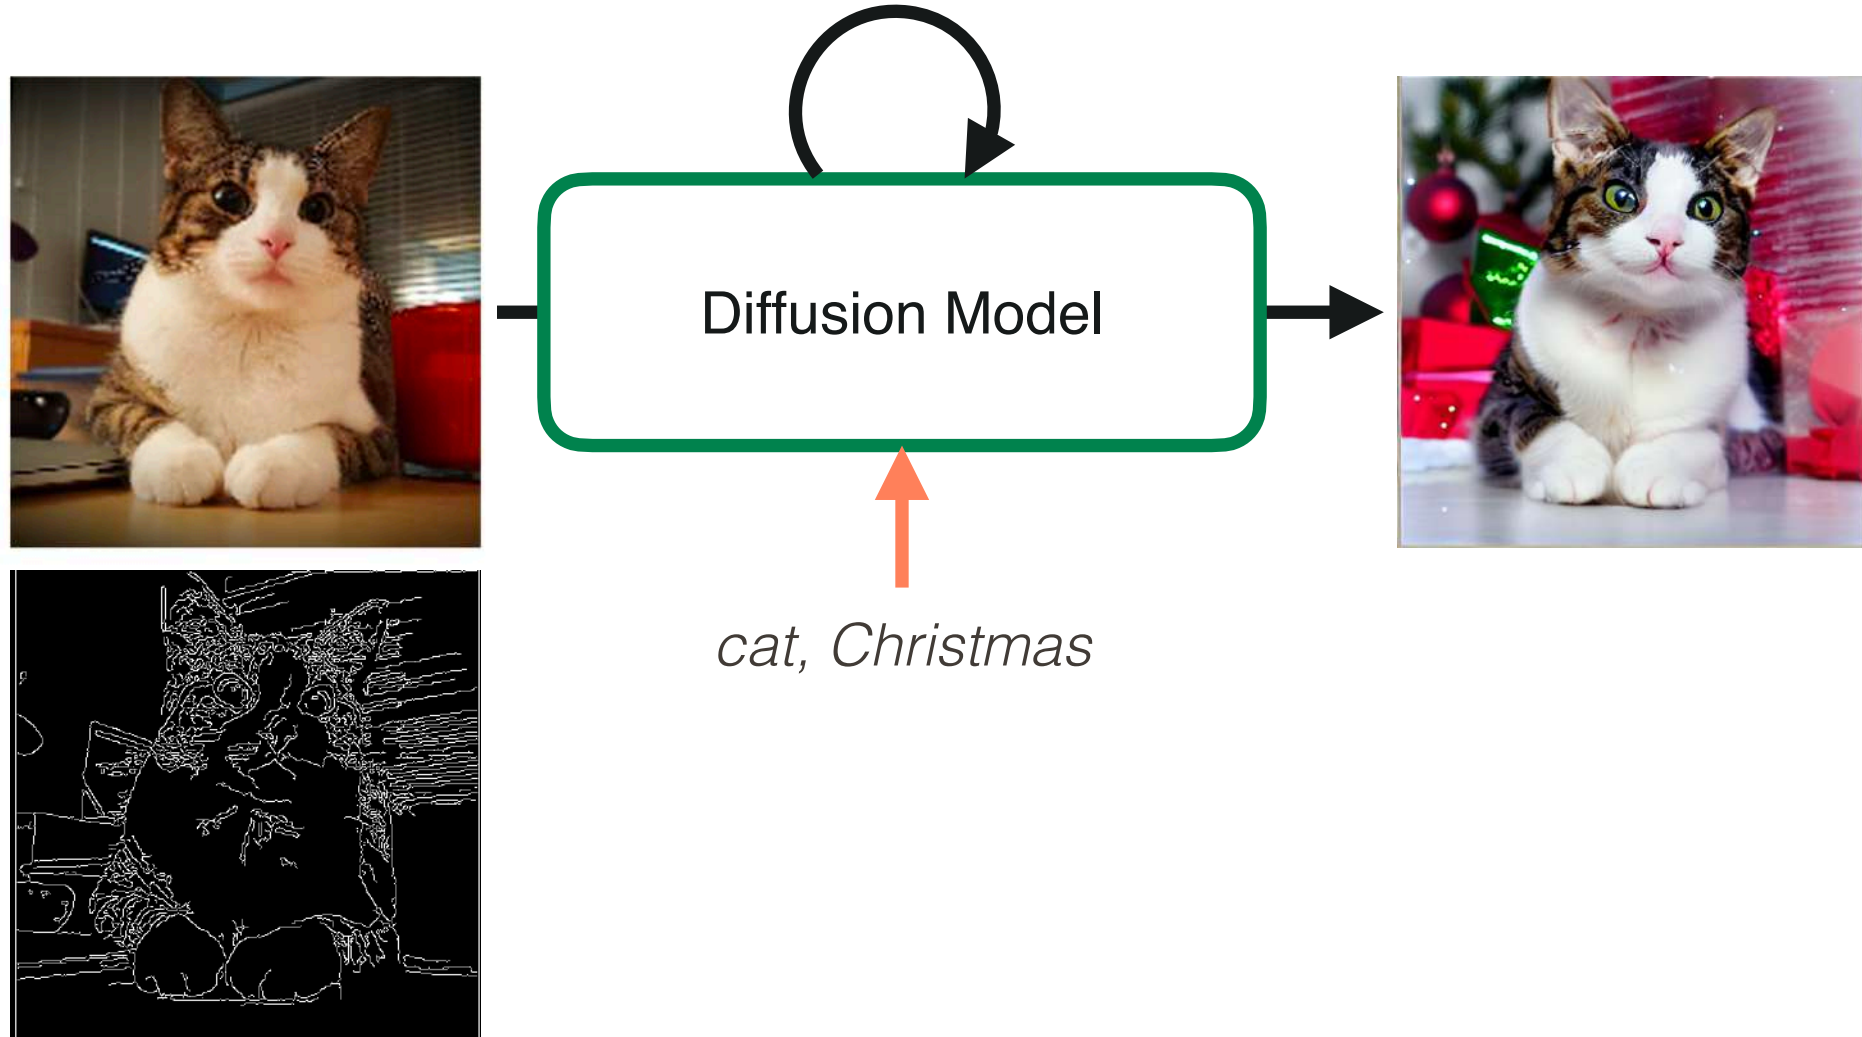


Image-to-Image Methods: Inpaint/Outpainting

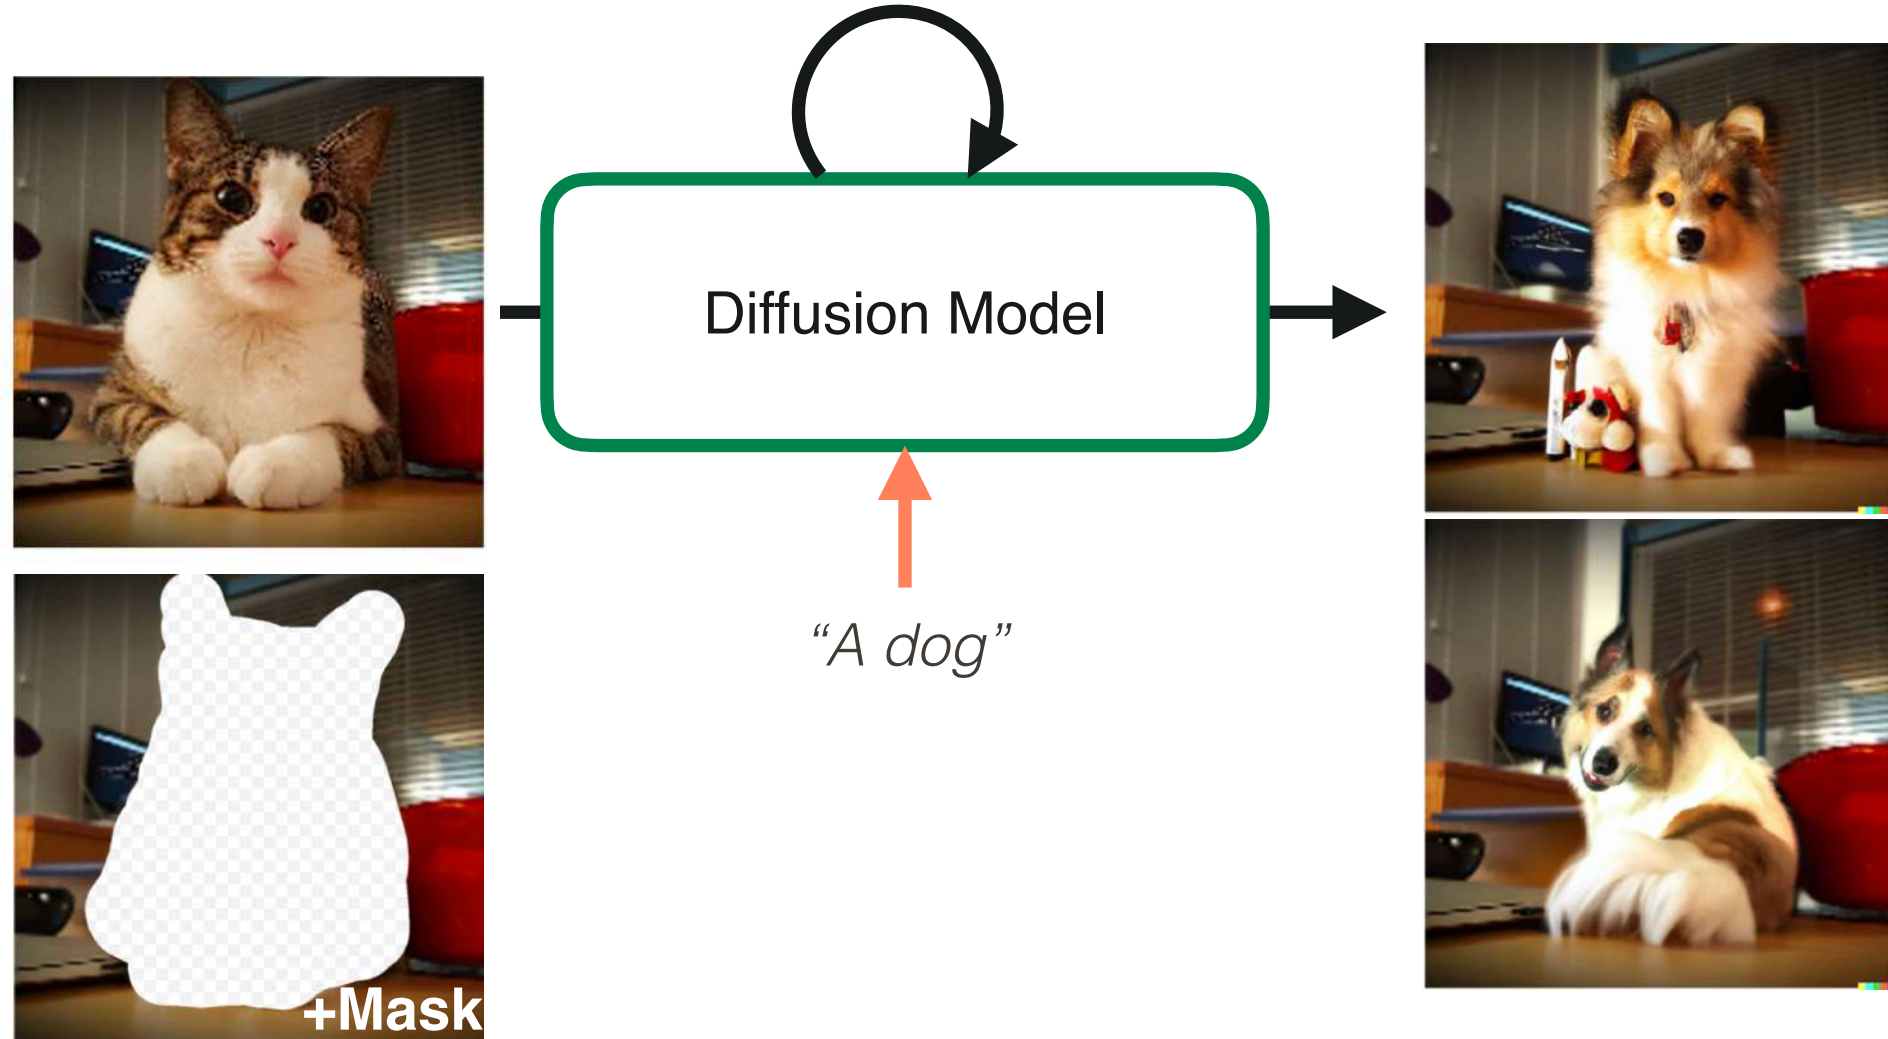
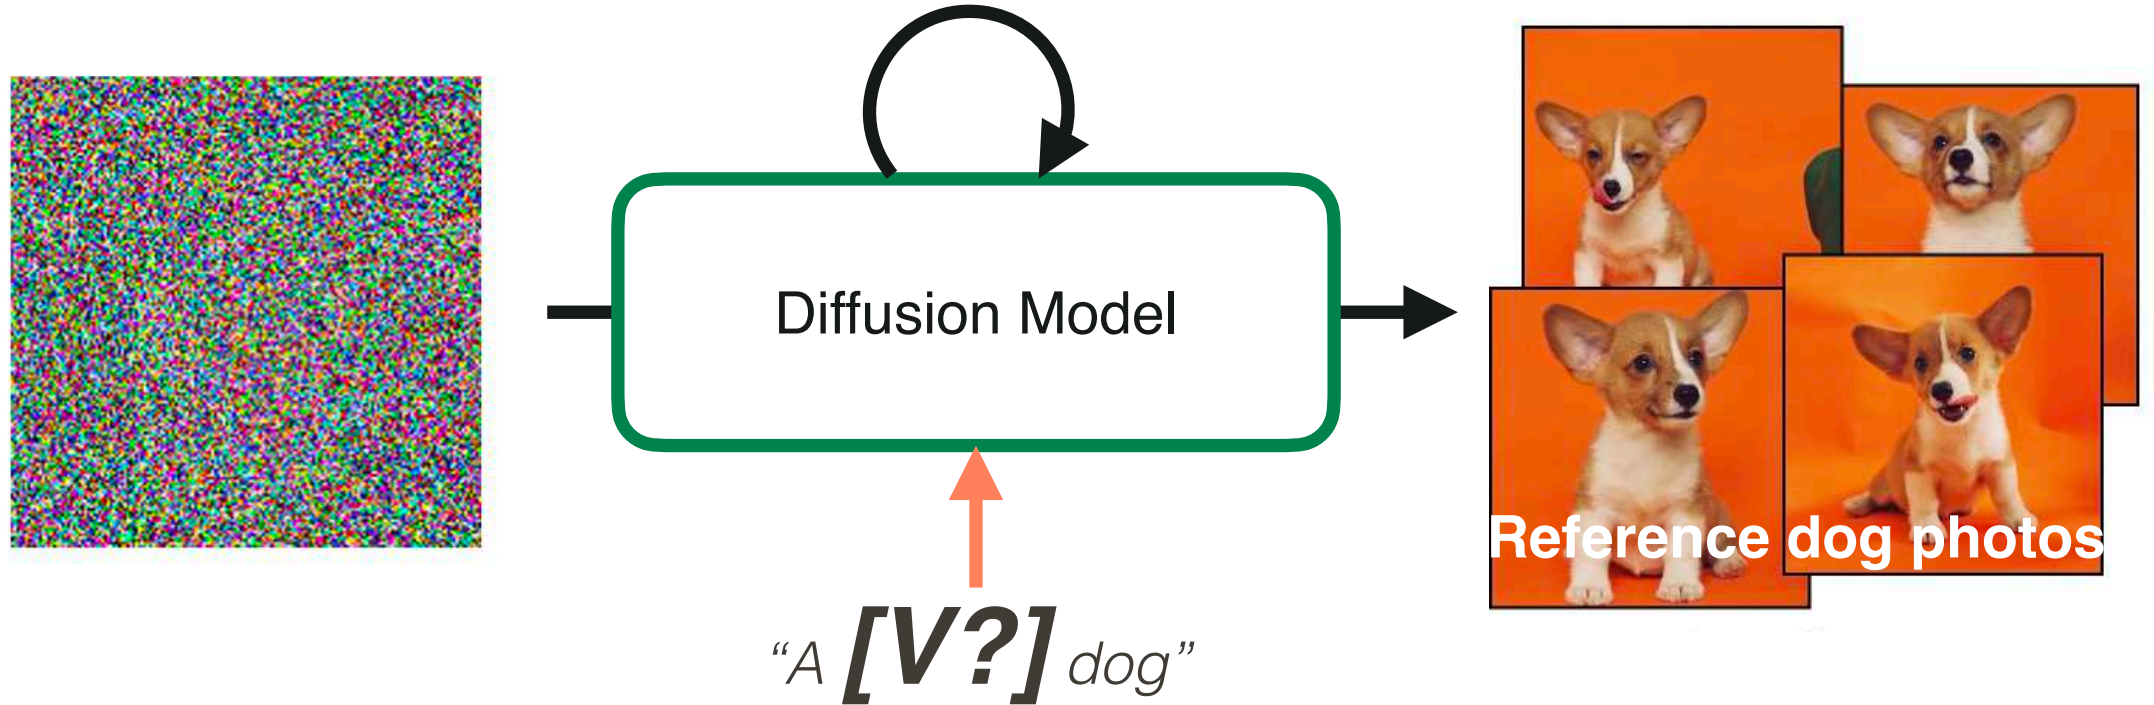
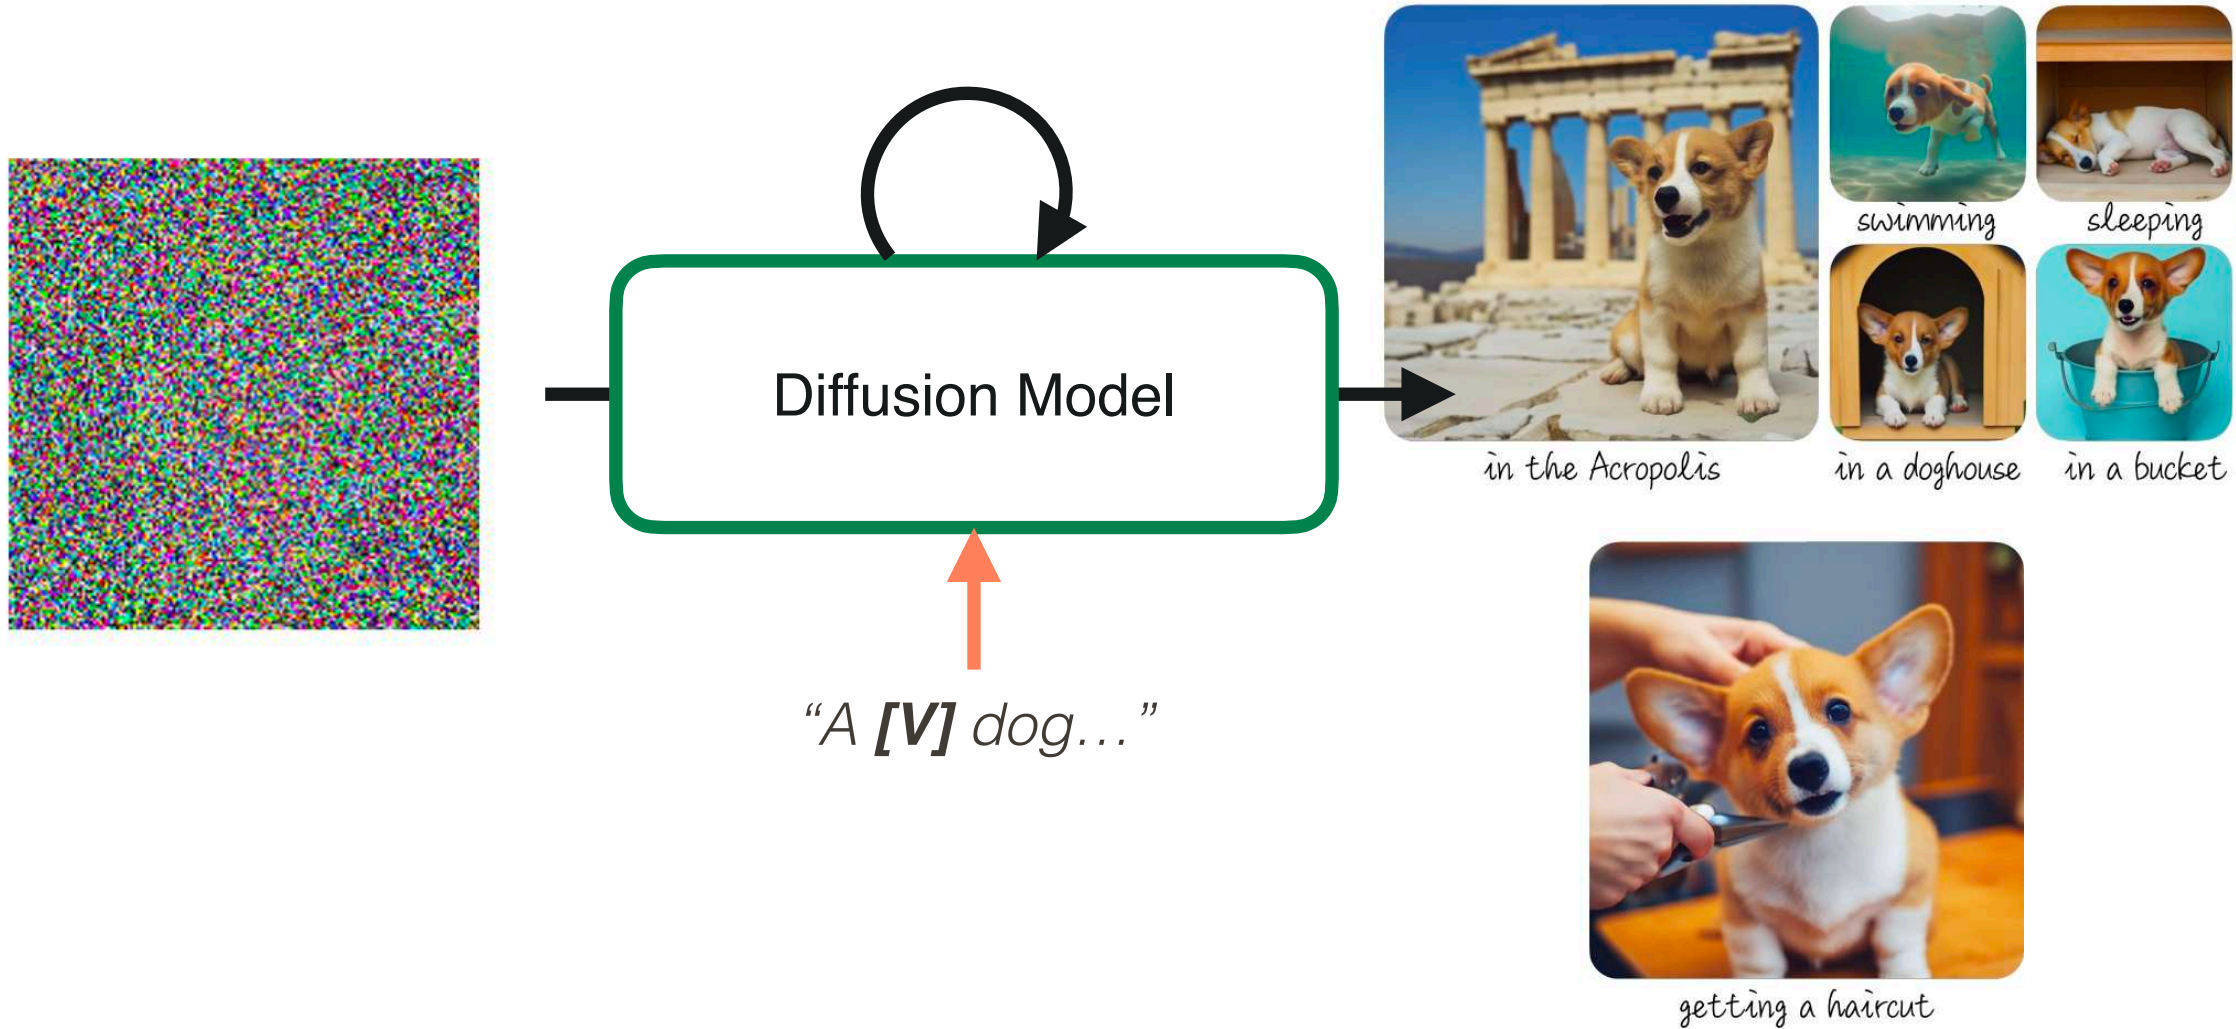


Image-to-Image Methods: Identity



* Based on Fig. 1, Ruiz et al., "DreamBooth: Fine Tuning Text-to-Image Diffusion Models for Subject-Driven Generation"

Image-to-Image Methods: Identity



* Based on Fig. 1, Ruiz et al., "DreamBooth: Fine Tuning Text-to-Image Diffusion Models for Subject-Driven Generation"

Strengths & Weaknesses

What it can do

1. Quick generation of complex, high-quality realistic and artistic images; good for creative exploration



From Aaron Hertzmann's blog


"cats in devo hats"

What it can do

2. Integrating specific styles



From Aaron Hertzmann's blog




"jacob lawrence painting of san francisco"

What it can do

2. Integrating specific styles



From Aaron Hertzmann's blog



"frank stella sculpture made of car parts"

What it can do

2. Integrating specific styles



"a monk riding a snail, medieval illuminated manuscript"

What it can't do

1. Following specific instructions (especially when the scene is complex):

- Composition



"a young dark-haired boy resting in bed, and a grey-haired older woman sitting in a chair beside the bed underneath a window with sun streaming through, Pixar style digital art"

What it can't do

1. Following specific instructions (especially when the scene is complex):

- Composition
- Generating multiple objects
- Coloring multiple objects

A yellow bowl and a blue cat



Catastrophic Neglect

One or more subjects are not generated

A yellow bowl and a brown bench



Incorrect Attribute Binding

Attributes (e.g., color) not matched correctly to subject

What it can't do

2. Being reasonably unbiased



"lawyer", April 6, 2022

"DALL·E 2 Preview - Risks and Limitations" by OpenAI

What it can't do

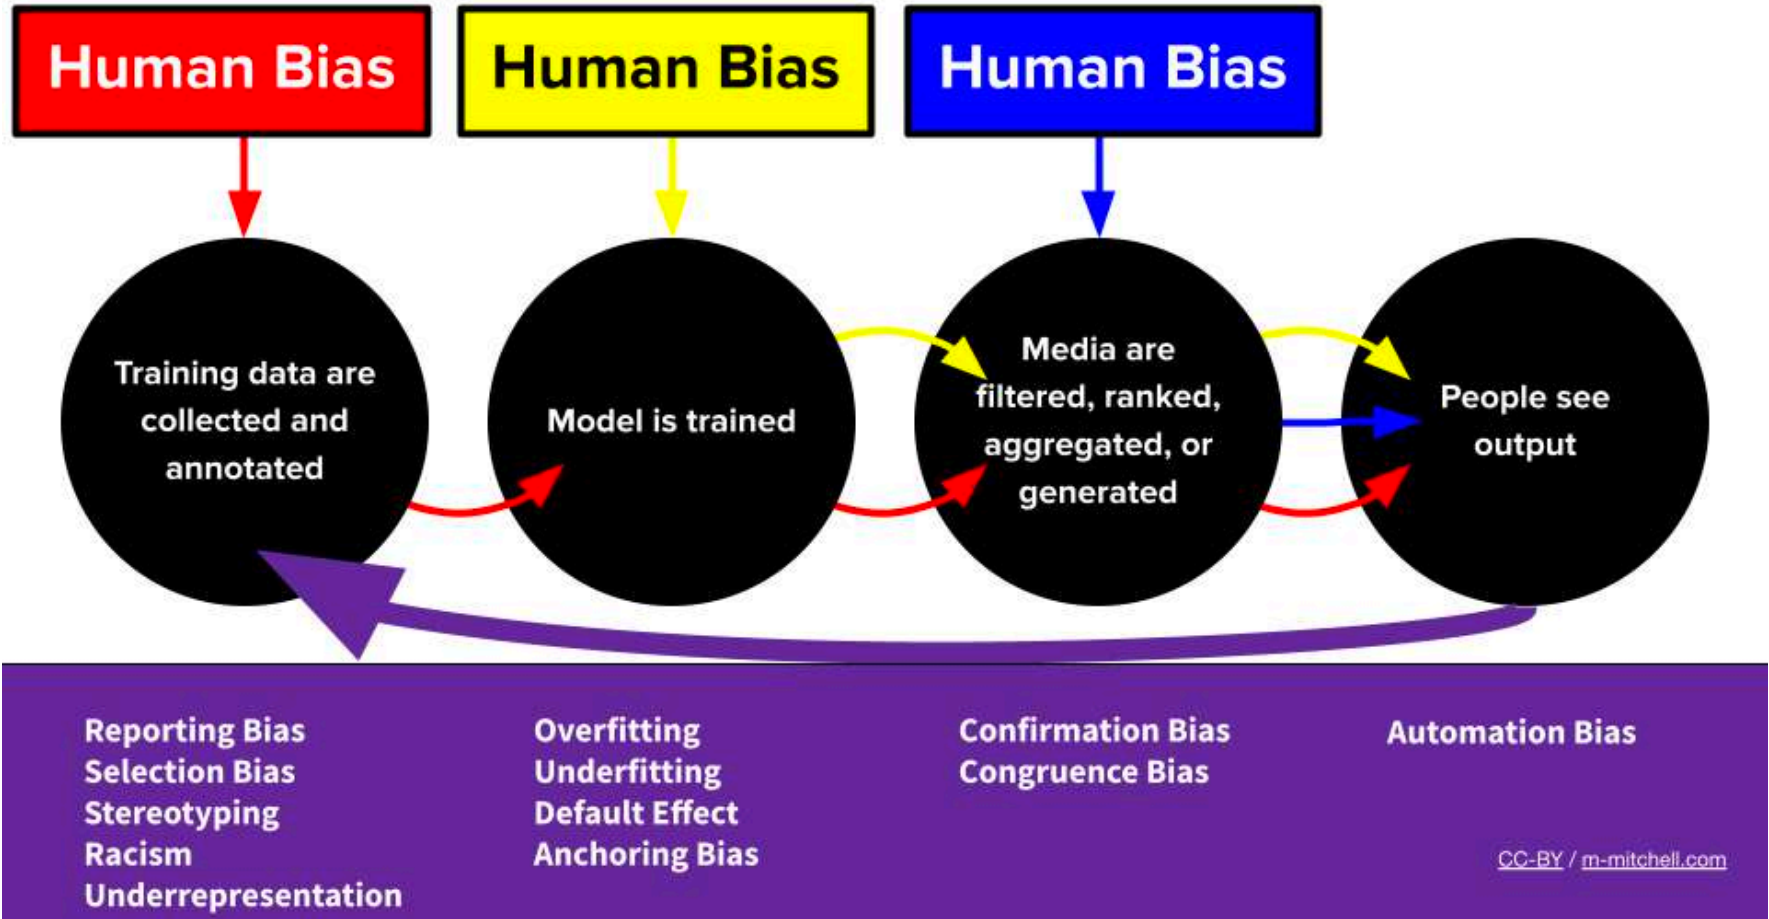
2. Being reasonably unbiased



“nurse”, April 6, 2022

“DALL·E 2 Preview - Risks and Limitations” by OpenAI

Where is bias from



The Bias ML Pipeline by Meg
<https://huggingface.co/meg>

How to be better

Bias can never be fully removed.

Heavily based on “Machine Learning in development:
Let's talk about bias!” By HuggingFace

How to be better: Addressing Bias

1. Task definition stage

- How ML techniques are integrated into the system? Is a ML model biased in a given use case?
- What is the optimization objective?

How to be better: Addressing Bias

2. Dataset selection and curation stage: A significant source of bias




Fig. 1: LAION5B training images: 10 most frequently occurring countries




Fig. 2: LAION5B most frequent job titles, showing unusually large number of monarchs

Cultural bias in AI models: a worked example with Barack Obama by Benjamin Peterson

How to be better: Addressing Bias

2. Dataset selection and curation stage

- Where is the data from? How was the dataset curated? What is the context?
- Measure the data. Any harmful associations?
- Document the dataset.
- Choose the dataset with least bias related harm. Iteratively improve the dataset.

How to be better: Addressing Bias

3. Model selection and training stage

- Visualize model outputs.
- Evaluate against benchmark.
- Document the model.

Generative Models & Artists

This is a fast evolving topic with many debates and open questions.

Warning: Contents may no longer be the state-of-art or relevant; and nothing presented should be taken as the “fact”.

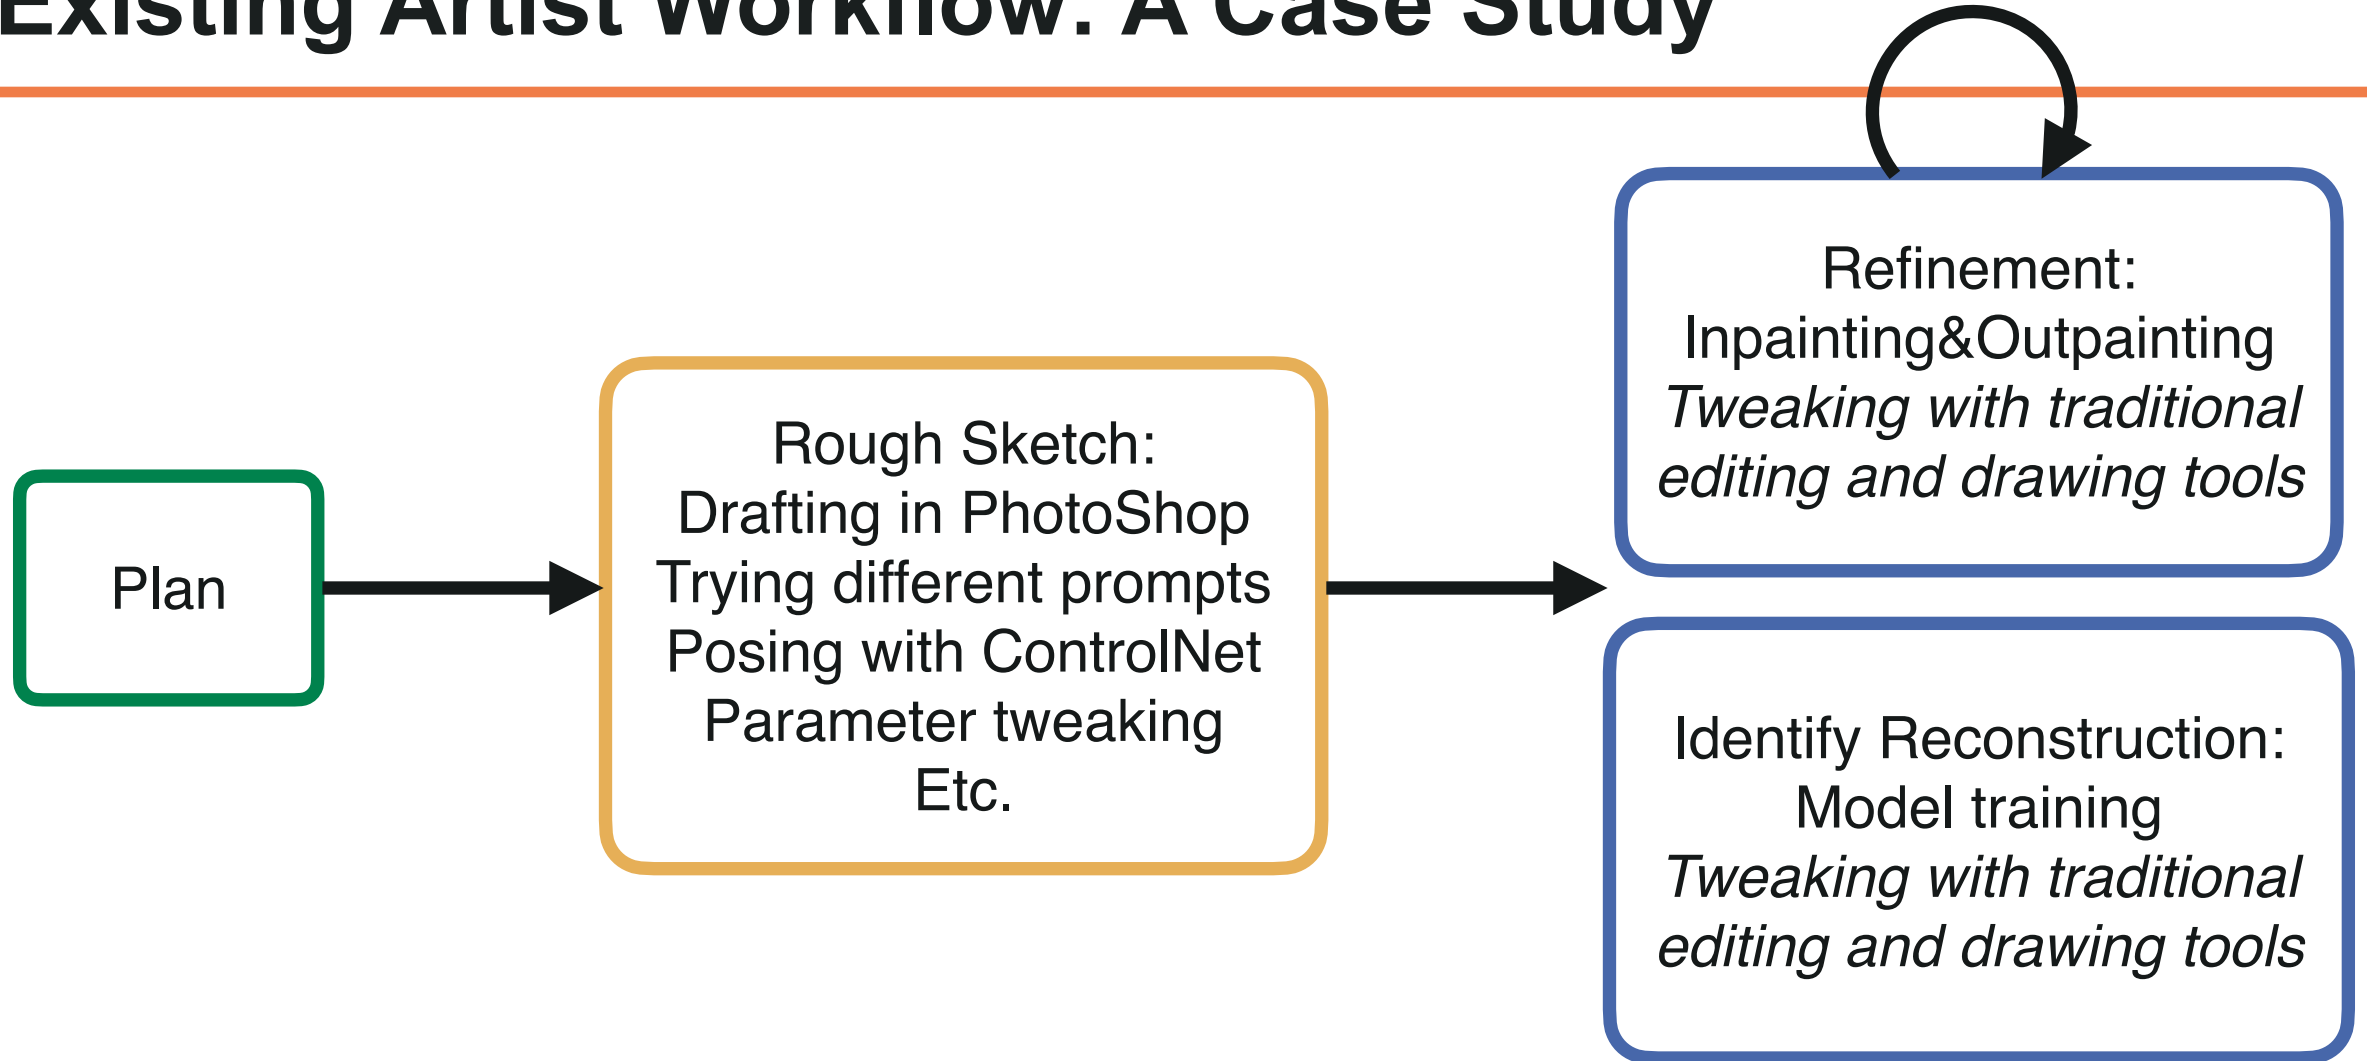
Existing Artist Workflow: A Case Study



"An AI artist explains his workflow"

<https://www.youtube.com/watch?v=K0ldxCh3cnI>

Existing Artist Workflow: A Case Study




"An AI artist explains his workflow"

<https://www.youtube.com/watch?v=K0ldxCh3cni>

Open Questions: How to protect artists?

Artist (V)


Original artwork



Glaze

Feature extractor (Φ)
Target style (T)

Cloaked artwork




Mimic



Cloaked artwork

fine-tune



Style-specific
model

generate



Fails to mimic
victim artist


Fig. 5, Shan et al., "Glaze: Protecting Artists from Style Mimicry by Text-to-Image Models"

Open Questions: How to attribute artists?

A significant concern of most participants, surprisingly, is not just the existence of AI art, but rather scraping of existing artworks without permission or compensation.

As one participant stated: “If artists are paid to have their pieces be used and asked permission, and if people had to pay to use that AI software with those pieces in it, I would have no problem.”

— Shan et al., “Glaze: Protecting Artists from Style Mimicry by Text-to-Image Models”



A model trained by an artist (DGSpitzer) on own drawings

Assignment Overview

Grading Policy

This assignment is graded subjectively.
We will be lenient.

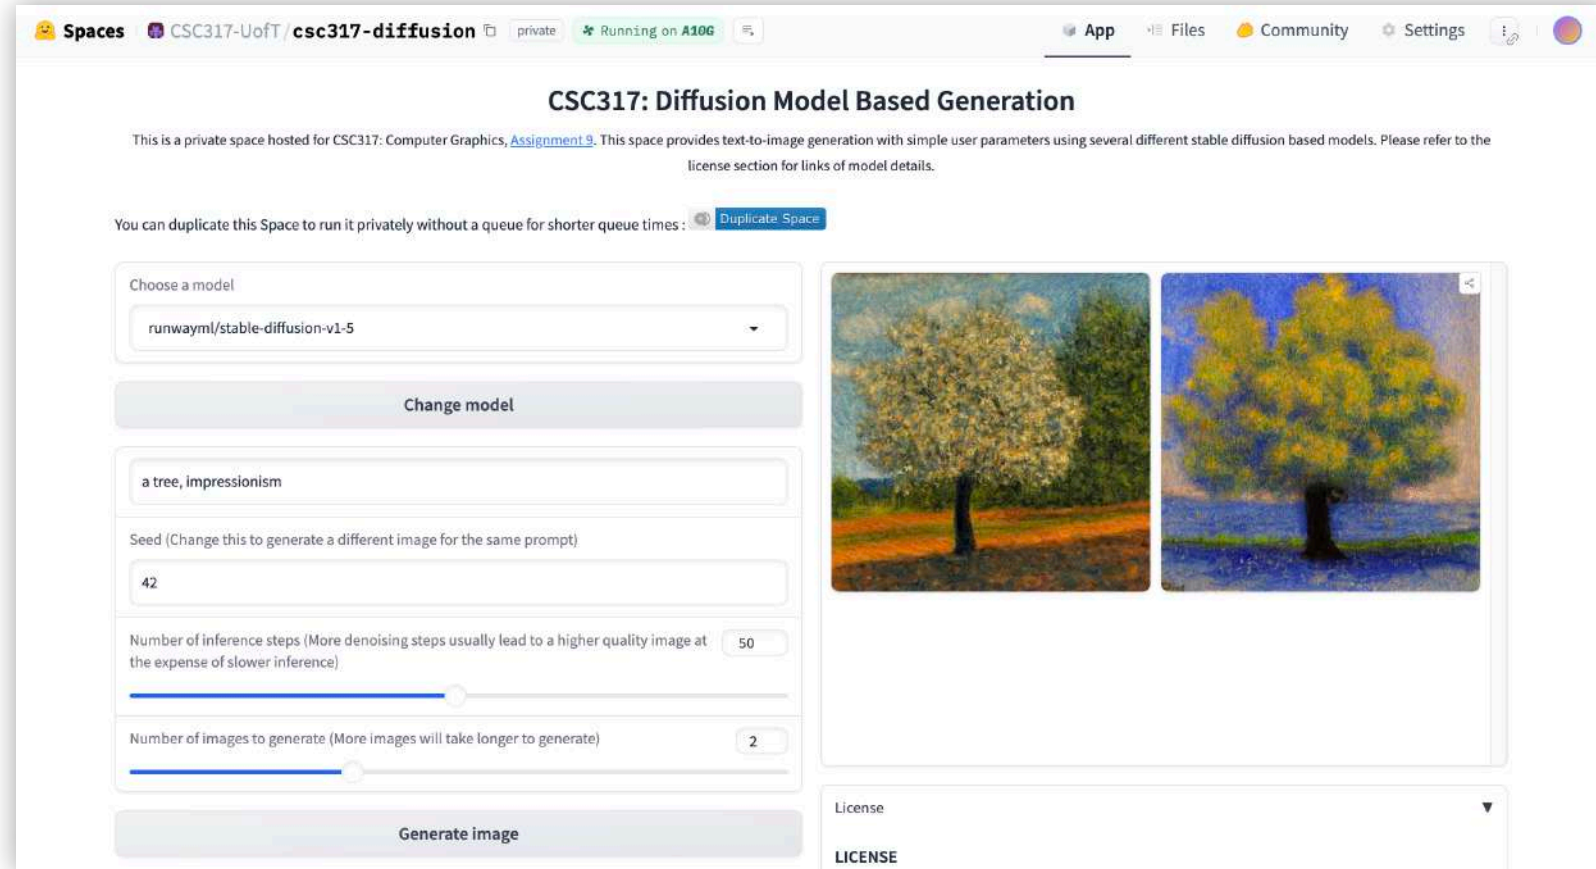
* You can request remarking if you question the mark

Privately-Hosted Generator


You will receive an invitation email by the end of today.

Contact us if:



1. You don't receive the email;
2. The queuing becomes too bad.
(We'll switch to better GPU before deadline)



Format



```
1  {
2    "pairs": [
3      {
4        "prompt": "[replace with prompt #1]",
5        "model": "[SDv1.5|SDv2.1|SDXL]",
6        "seed": 42,
7        "image-path": "image-1.jpg"
8      },
9      {
10       "prompt": "[replace with prompt #2]",
11       "model": "[SDv1.5|SDv2.1|SDXL]",
12       "seed": 42,
13       "image-path": "image-2.jpg"
14     },
15     {
16       "prompt": "[replace with prompt #3]",
17       "model": "[SDv1.5|SDv2.1|SDXL]",
18       "seed": 42,
19       "image-path": "image-3.jpg"
20     },
21     {
22       "prompt": "[replace with prompt #4]",
23       "model": "[SDv1.5|SDv2.1|SDXL]",
24       "seed": 42,
25       "image-path": "image-4.jpg"
26     },
27     {
28       "prompt": "[replace with prompt #5]",
29       "model": "[SDv1.5|SDv2.1|SDXL]",
30       "seed": 42,
31       "image-path": "image-5.jpg"
32     },
33     {
34       "prompt": "[replace with prompt #6]",
35       "model": "[SDv1.5|SDv2.1|SDXL]",
36       "seed": 42,
37       "image-path": "image-6.jpg"
38     }
39   ],
40   "story": "[replace with description of the story taking place.]"
41 }
```



[FEED ME]

Drag & drop
either all files or a .zip
for a specific task

Be Reasonable

Only use the generator for this assignment.
Only submit images generated by our setup.

Awards

We'll pick and frame THREE "open-ended" or "story" images



Awards

*We'll fund the author of the best image to
SIGGRAPH next year!*




SIGGRAPH 2024

DENVER+ 28 JUL — 1 AUG

The 51st International Conference & Exhibition On
Computer Graphics & Interactive Techniques



Sponsored by ACM SIGGRAPH



Thank you!
Questions?

Further Readings

Intro

- Zhu, Xiaojin, et al. "A text-to-picture synthesis system for augmenting communication." *AAAI*. Vol. 7. 2007.
<https://pages.cs.wisc.edu/~jerryzhu/pub/ttp.pdf>
- Mansimov, Elman, et al. "Generating images from captions with attention." *arXiv preprint arXiv:1511.02793* (2015).
<https://arxiv.org/pdf/1511.02793.pdf>
- Ho, Jonathan, Ajay Jain, and Pieter Abbeel. "Denoising diffusion probabilistic models." *Advances in neural information processing systems* 33 (2020): 6840-6851.
<https://hojonathanho.github.io/diffusion/>
- Dhariwal, Prafulla, and Alexander Nichol. "Diffusion models beat gans on image synthesis." *Advances in neural information processing systems* 34 (2021): 8780-8794.
<https://arxiv.org/abs/2105.05233>
- Reed, Scott, et al. "Generative adversarial text to image synthesis." *International conference on machine learning*. PMLR, 2016.
<https://proceedings.mlr.press/v48/reed16.pdf>

A Handwavy Introduction to Diffusion Model

- Weng, Lilian. (Jul 2021). What are diffusion models? Lil'Log.
<https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>
- Song, Yang. (May 2021). Generative Modeling by Estimating Gradients of the Data Distribution. Yang Song's blog.
<https://yang-song.net/blog/2021/score/>
- Yang Song's tutorial video.
<https://www.youtube.com/watch?v=wMmqCMwuM2Q>

A Handwavy Introduction to Diffusion Model

- Vaclav Kosar. Cross-Attention in Transformer Architecture.
<https://vaclavkosar.com/ml/cross-attention-in-transformer-architecture>
- Tang, Raphael, et al. "What the daam: Interpreting stable diffusion using cross attention." arXiv preprint arXiv:2210.04885 (2022).
<https://github.com/castorini/daam>
- Hertz, Amir, et al. "Prompt-to-prompt image editing with cross attention control." arXiv preprint arXiv:2208.01626 (2022).
<https://prompt-to-prompt.github.io/>

Advanced Diffusion-Model-Based Editing Tools

- Zhang, Lvmin, Anyi Rao, and Maneesh Agrawala. "Adding conditional control to text-to-image diffusion models." *Proceedings of the IEEE/CVF International Conference on Computer Vision*. 2023.

<https://github.com/lllyasviel/ControlNet>

- Ruiz, Nataniel, et al. "Dreambooth: Fine tuning text-to-image diffusion models for subject-driven generation." *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2023.

<https://dreambooth.github.io/>

Strengths & Weaknesses

- Aaron Hertzmann's blog. Creative Explorations with DALL-E 2.
<https://aaronhertzmann.com/2022/05/25/dall-e.html>
- Miranda Dixon-Luinenburg. What DALL-E 2 can and cannot do.
<https://www.lesswrong.com/posts/uKp6tBFStnsvrot5t/what-dall-e-2-can-and-cannot-do>
- Chefer, Hila, et al. "Attend-and-excite: Attention-based semantic guidance for text-to-image diffusion models." ACM Transactions on Graphics (TOG) 42.4 (2023): 1-10.
<https://github.com/yuval-alaluf/Attend-and-Excite>

Strengths & Weaknesses

- OpenAI. (Jul 2022). DALL·E 2 Preview - Risks and Limitations. <https://github.com/openai/dalle-2-preview/blob/main/system-card.md#bias-and-representation>
- Jernite, Yacine, et al. (Aug 2023). Hugging Face Ethics and Society Newsletter 2: Let's Talk about Bias!. Hugging Face Blog. <https://huggingface.co/blog/ethics-soc-2>

Generative Models & Artists

- “An AI artist explains his workflow”

<https://www.youtube.com/watch?v=K0ldxCh3cnI>

- Shan, Shawn, et al. "Glaze: Protecting artists from style mimicry by text-to-image models." arXiv preprint arXiv:2302.04222 (2023).

<https://arxiv.org/abs/2302.04222>

- DGSpitzer Art Diffusion.

<https://huggingface.co/DGSpitzer/DGSpitzer-Art-Diffusion>